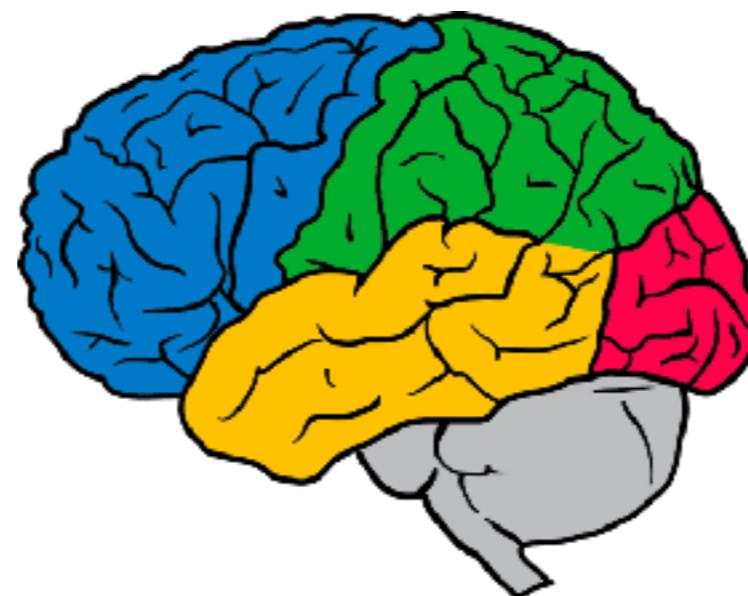


Adversarial Examples and Adversarial Training

Ian Goodfellow, Staff Research Scientist, Google Brain

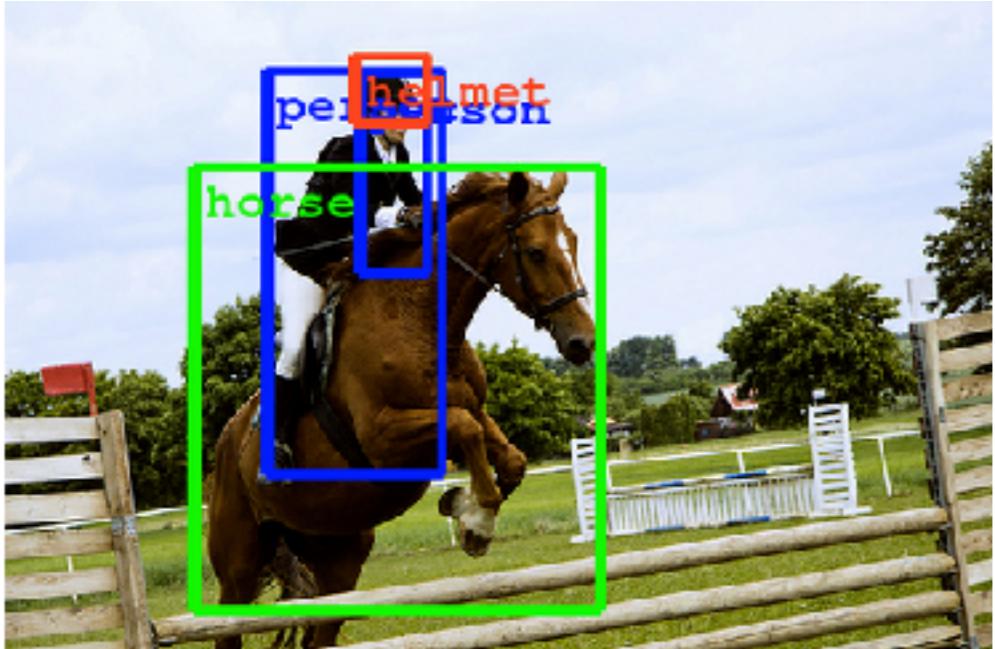
CS 231n, Stanford University, 2017-05-30



Overview

- What are adversarial examples?
- Why do they happen?
- How can they be used to compromise machine learning systems?
- What are the defenses?
- How to use adversarial examples to improve machine learning, even when there is no adversary

Since 2013, deep neural networks have matched human performance at...



(Szegedy et al, 2014)

...recognizing objects
and faces....



(Taigmen et al, 2013)



(Goodfellow et al, 2013)

...solving CAPTCHAS and
reading addresses...



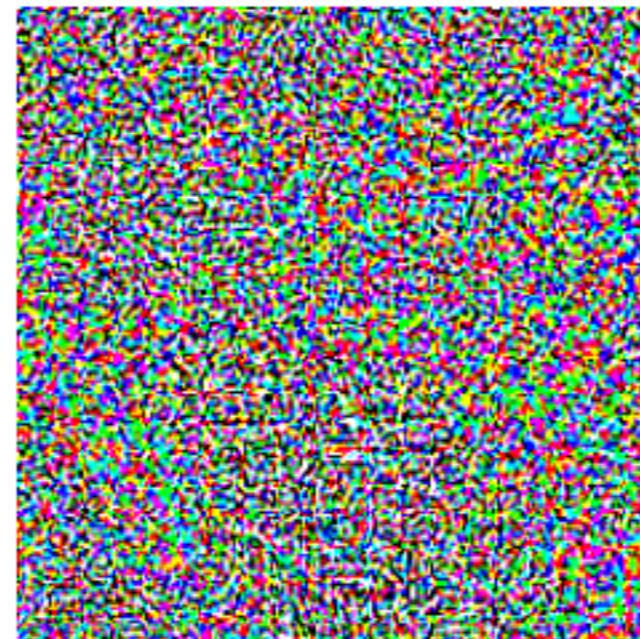
(Goodfellow et al, 2013)

and other tasks...

Adversarial Examples



+ .007 ×



=



Timeline:

“Adversarial Classification” Dalvi et al 2004: fool spam filter

“Evasion Attacks Against Machine Learning at Test Time”

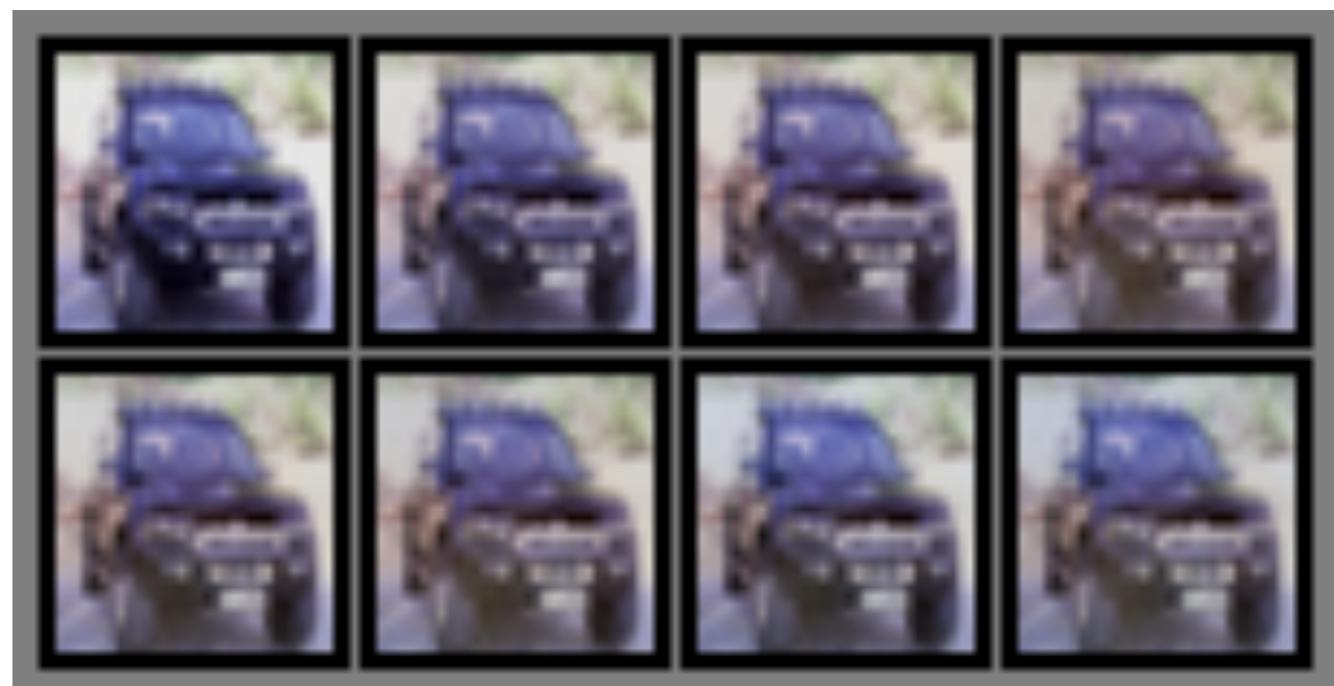
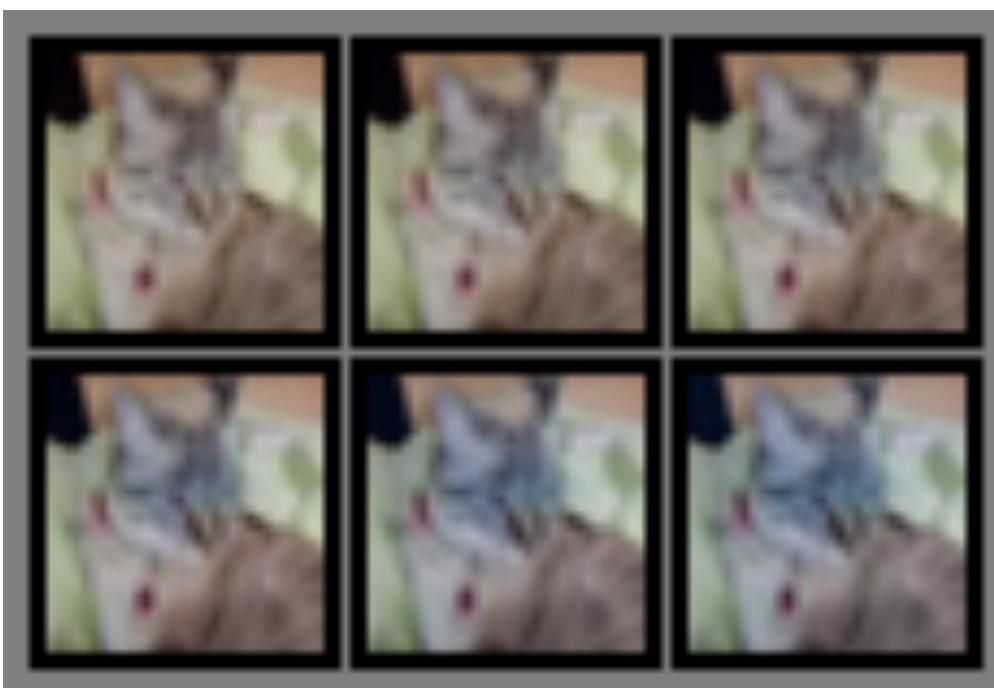
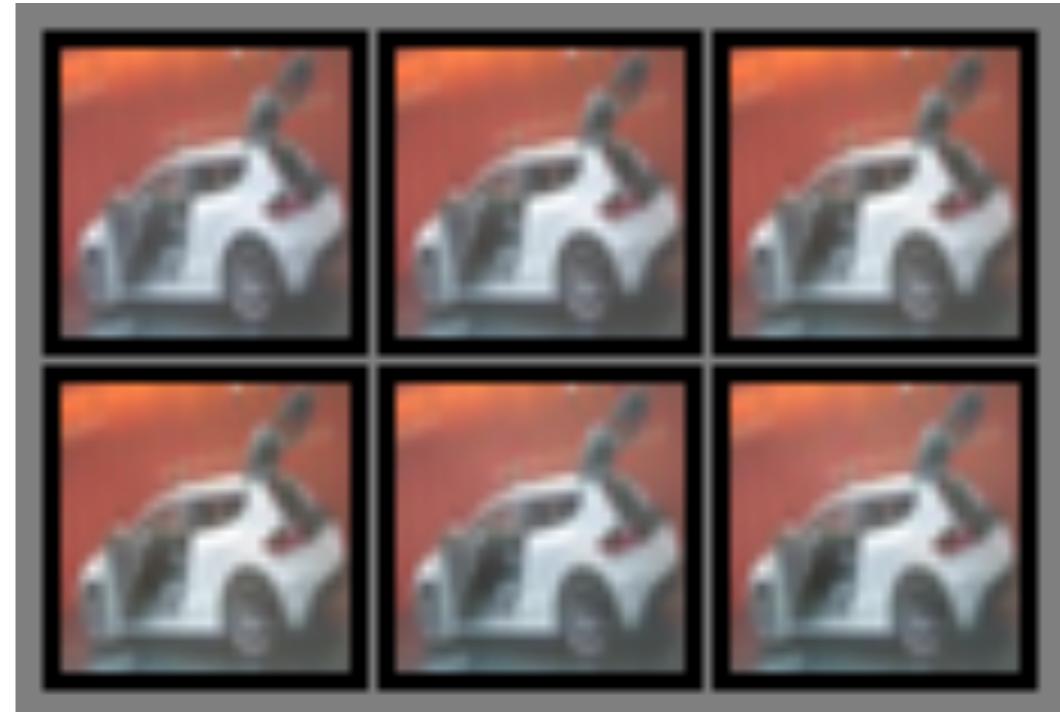
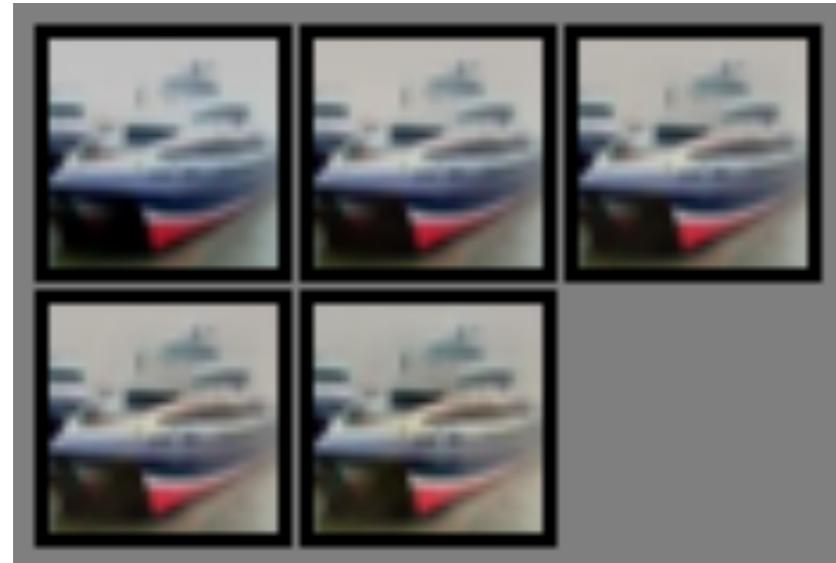
Biggio 2013: fool neural nets

Szegedy et al 2013: fool ImageNet classifiers imperceptibly

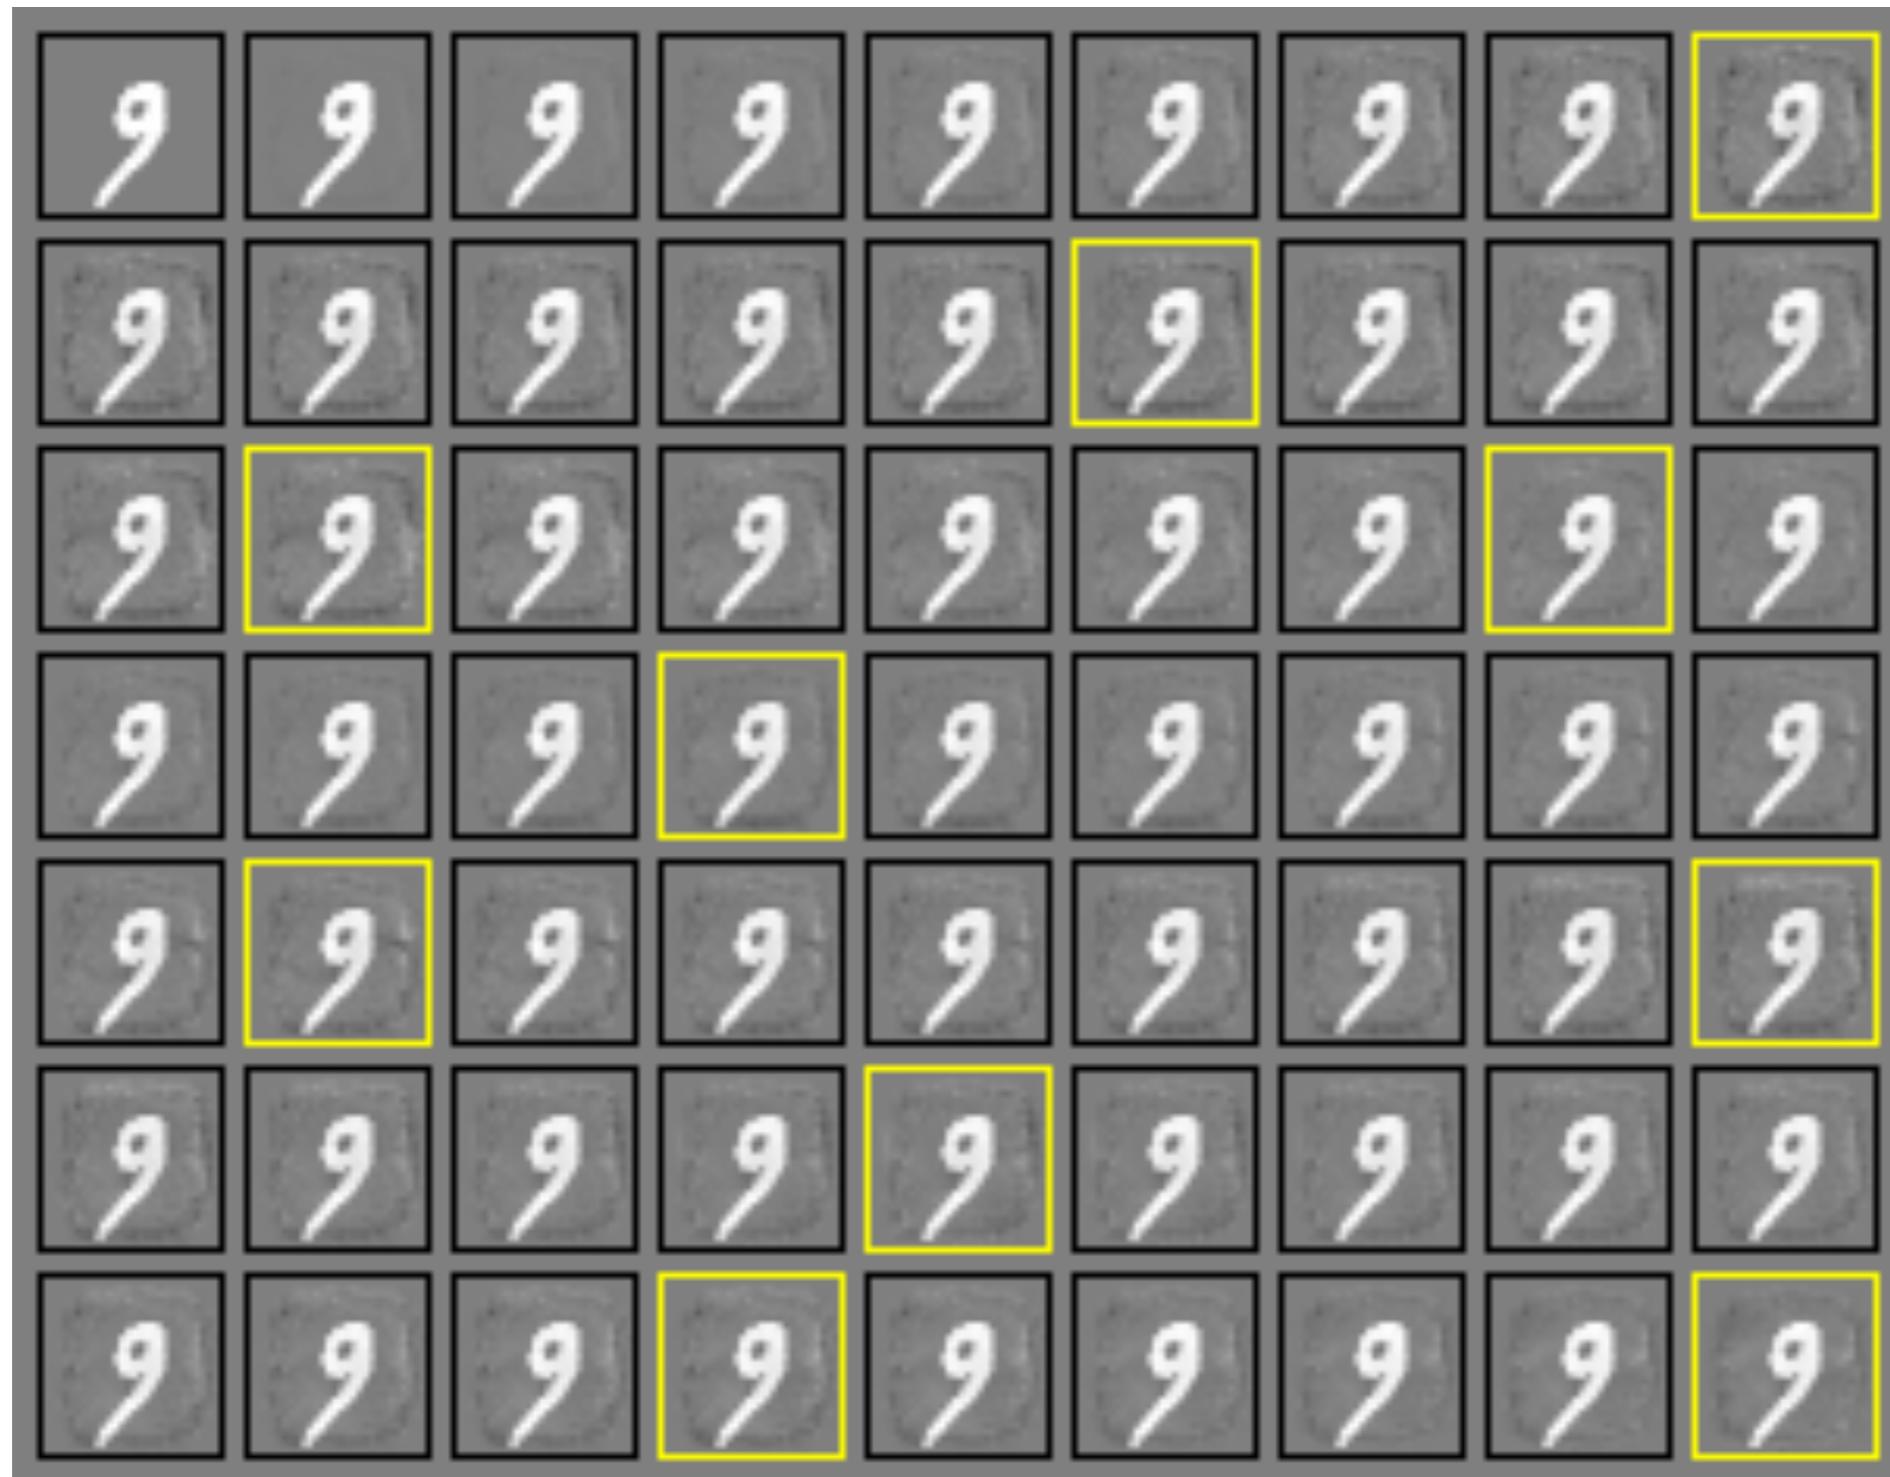
Goodfellow et al 2014: cheap, closed form attack

(Goodfellow 2016)

Turning Objects into “Airplanes”



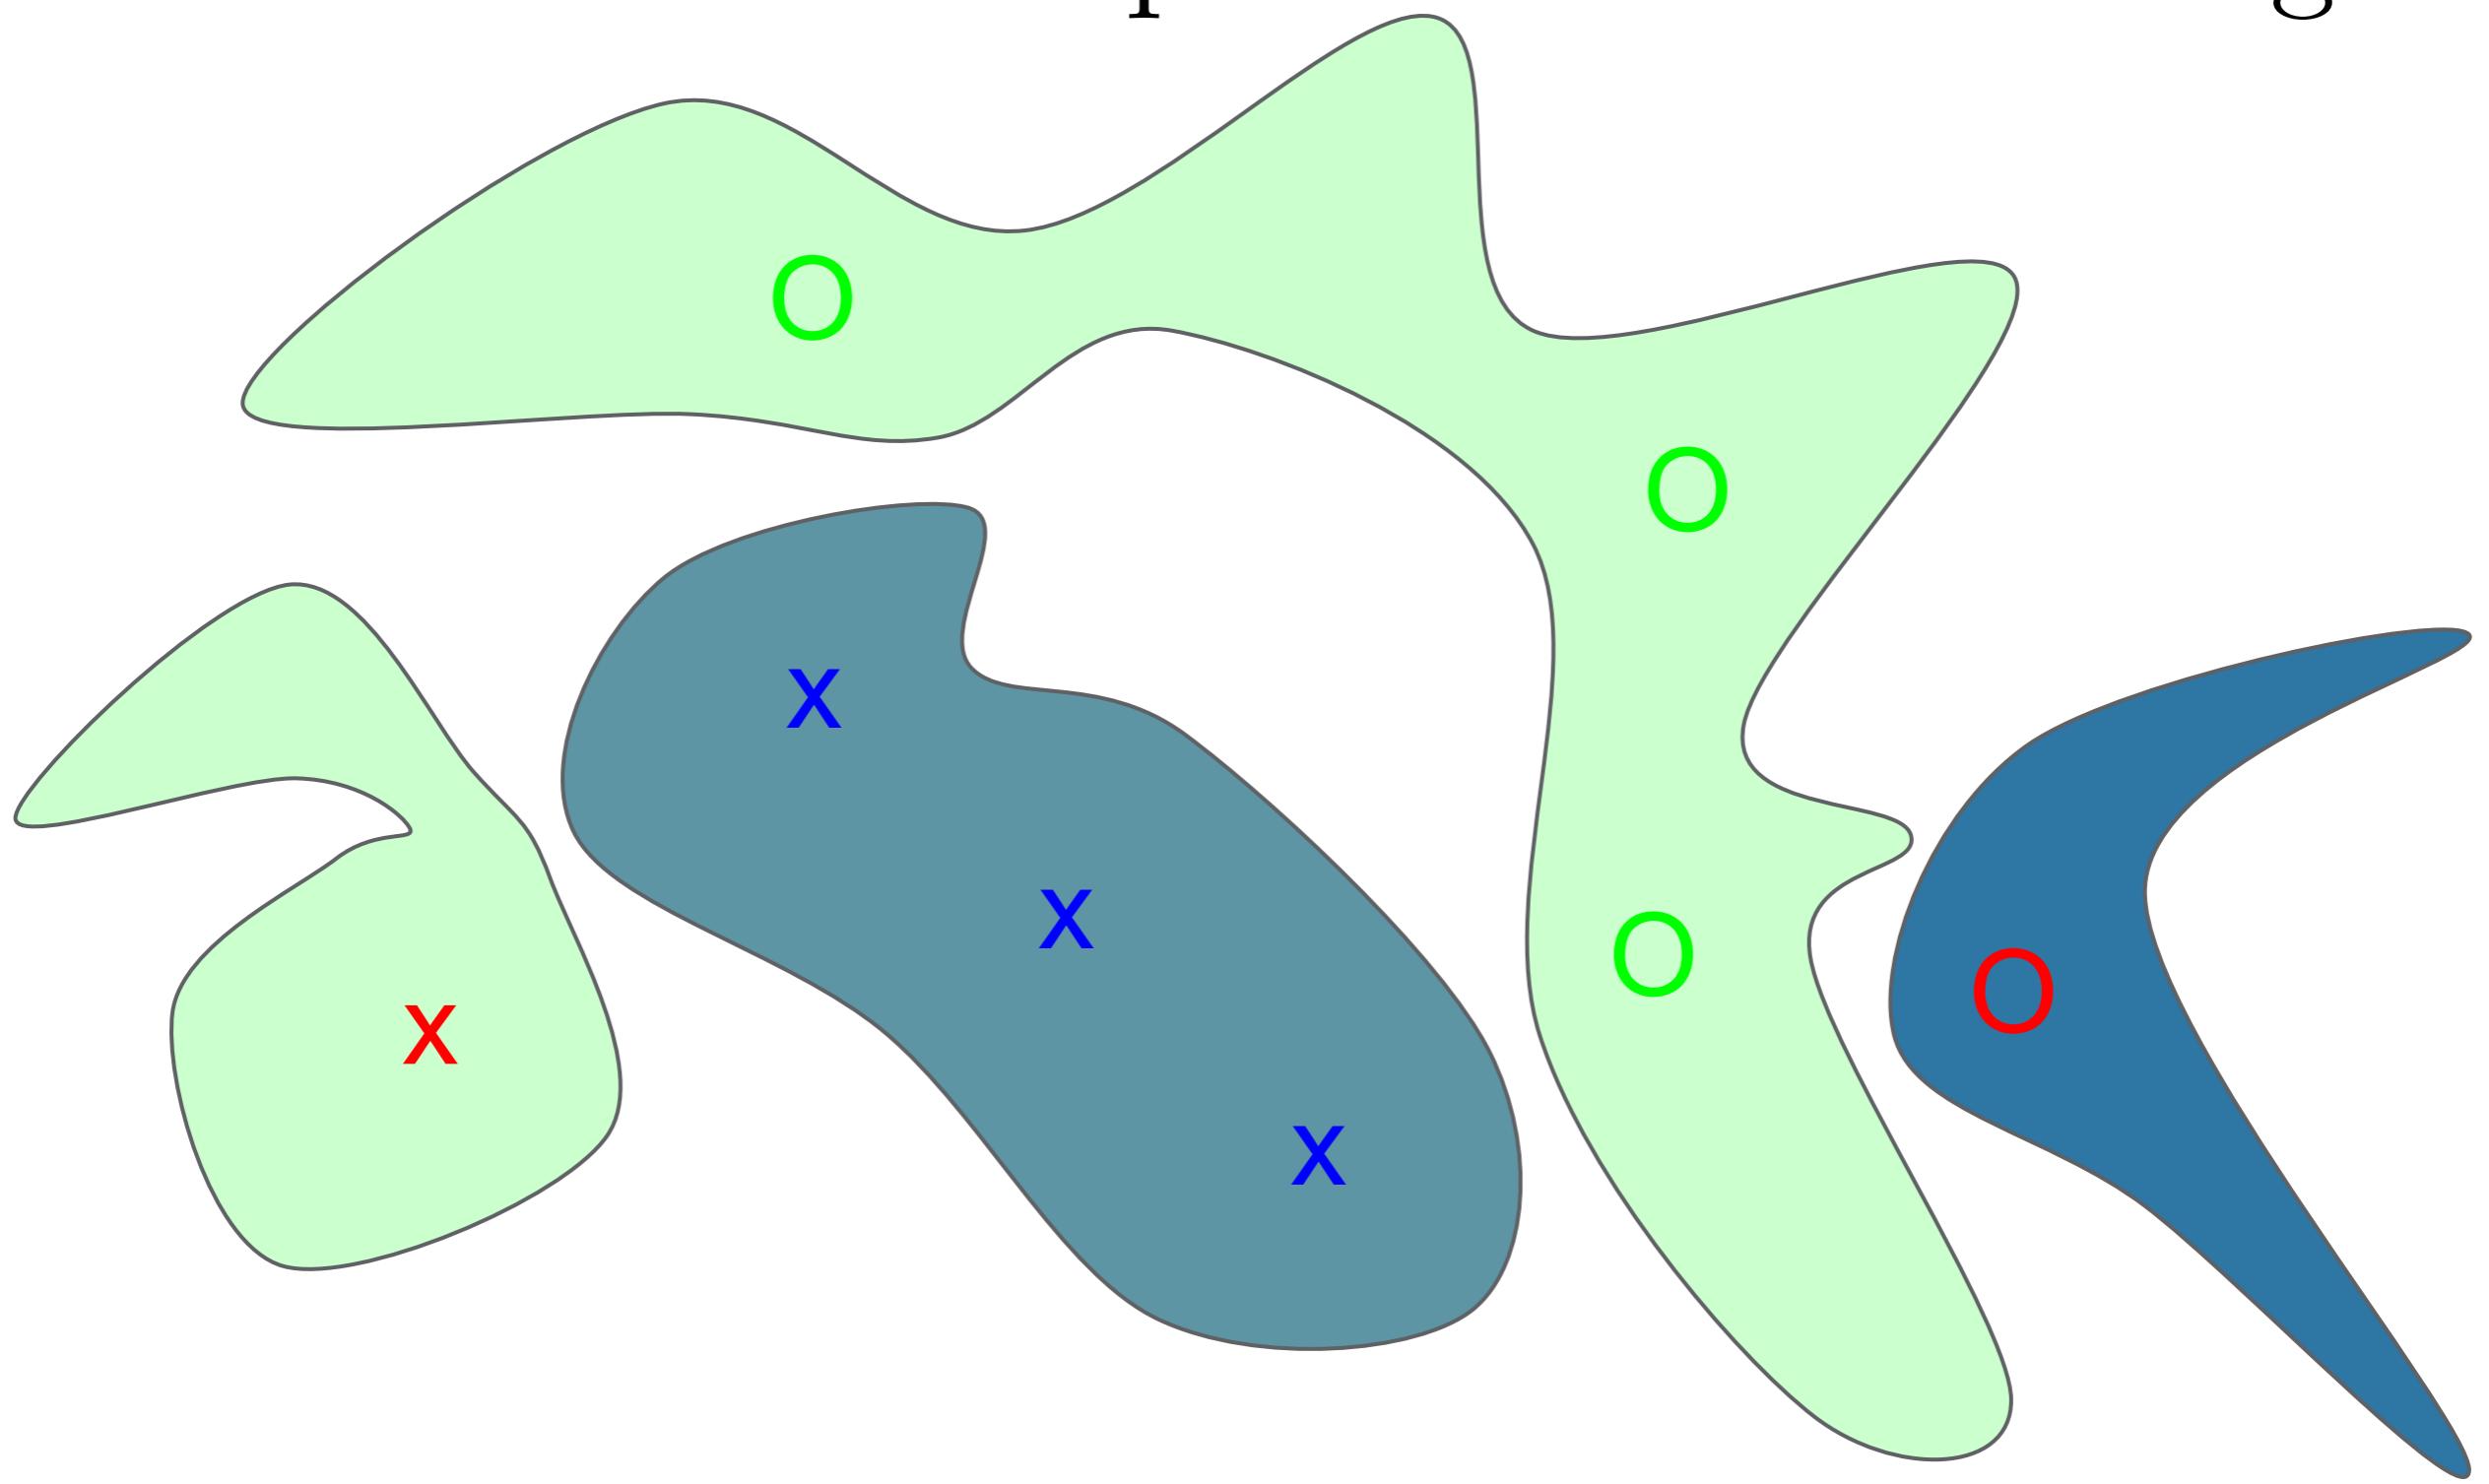
Attacking a Linear Model



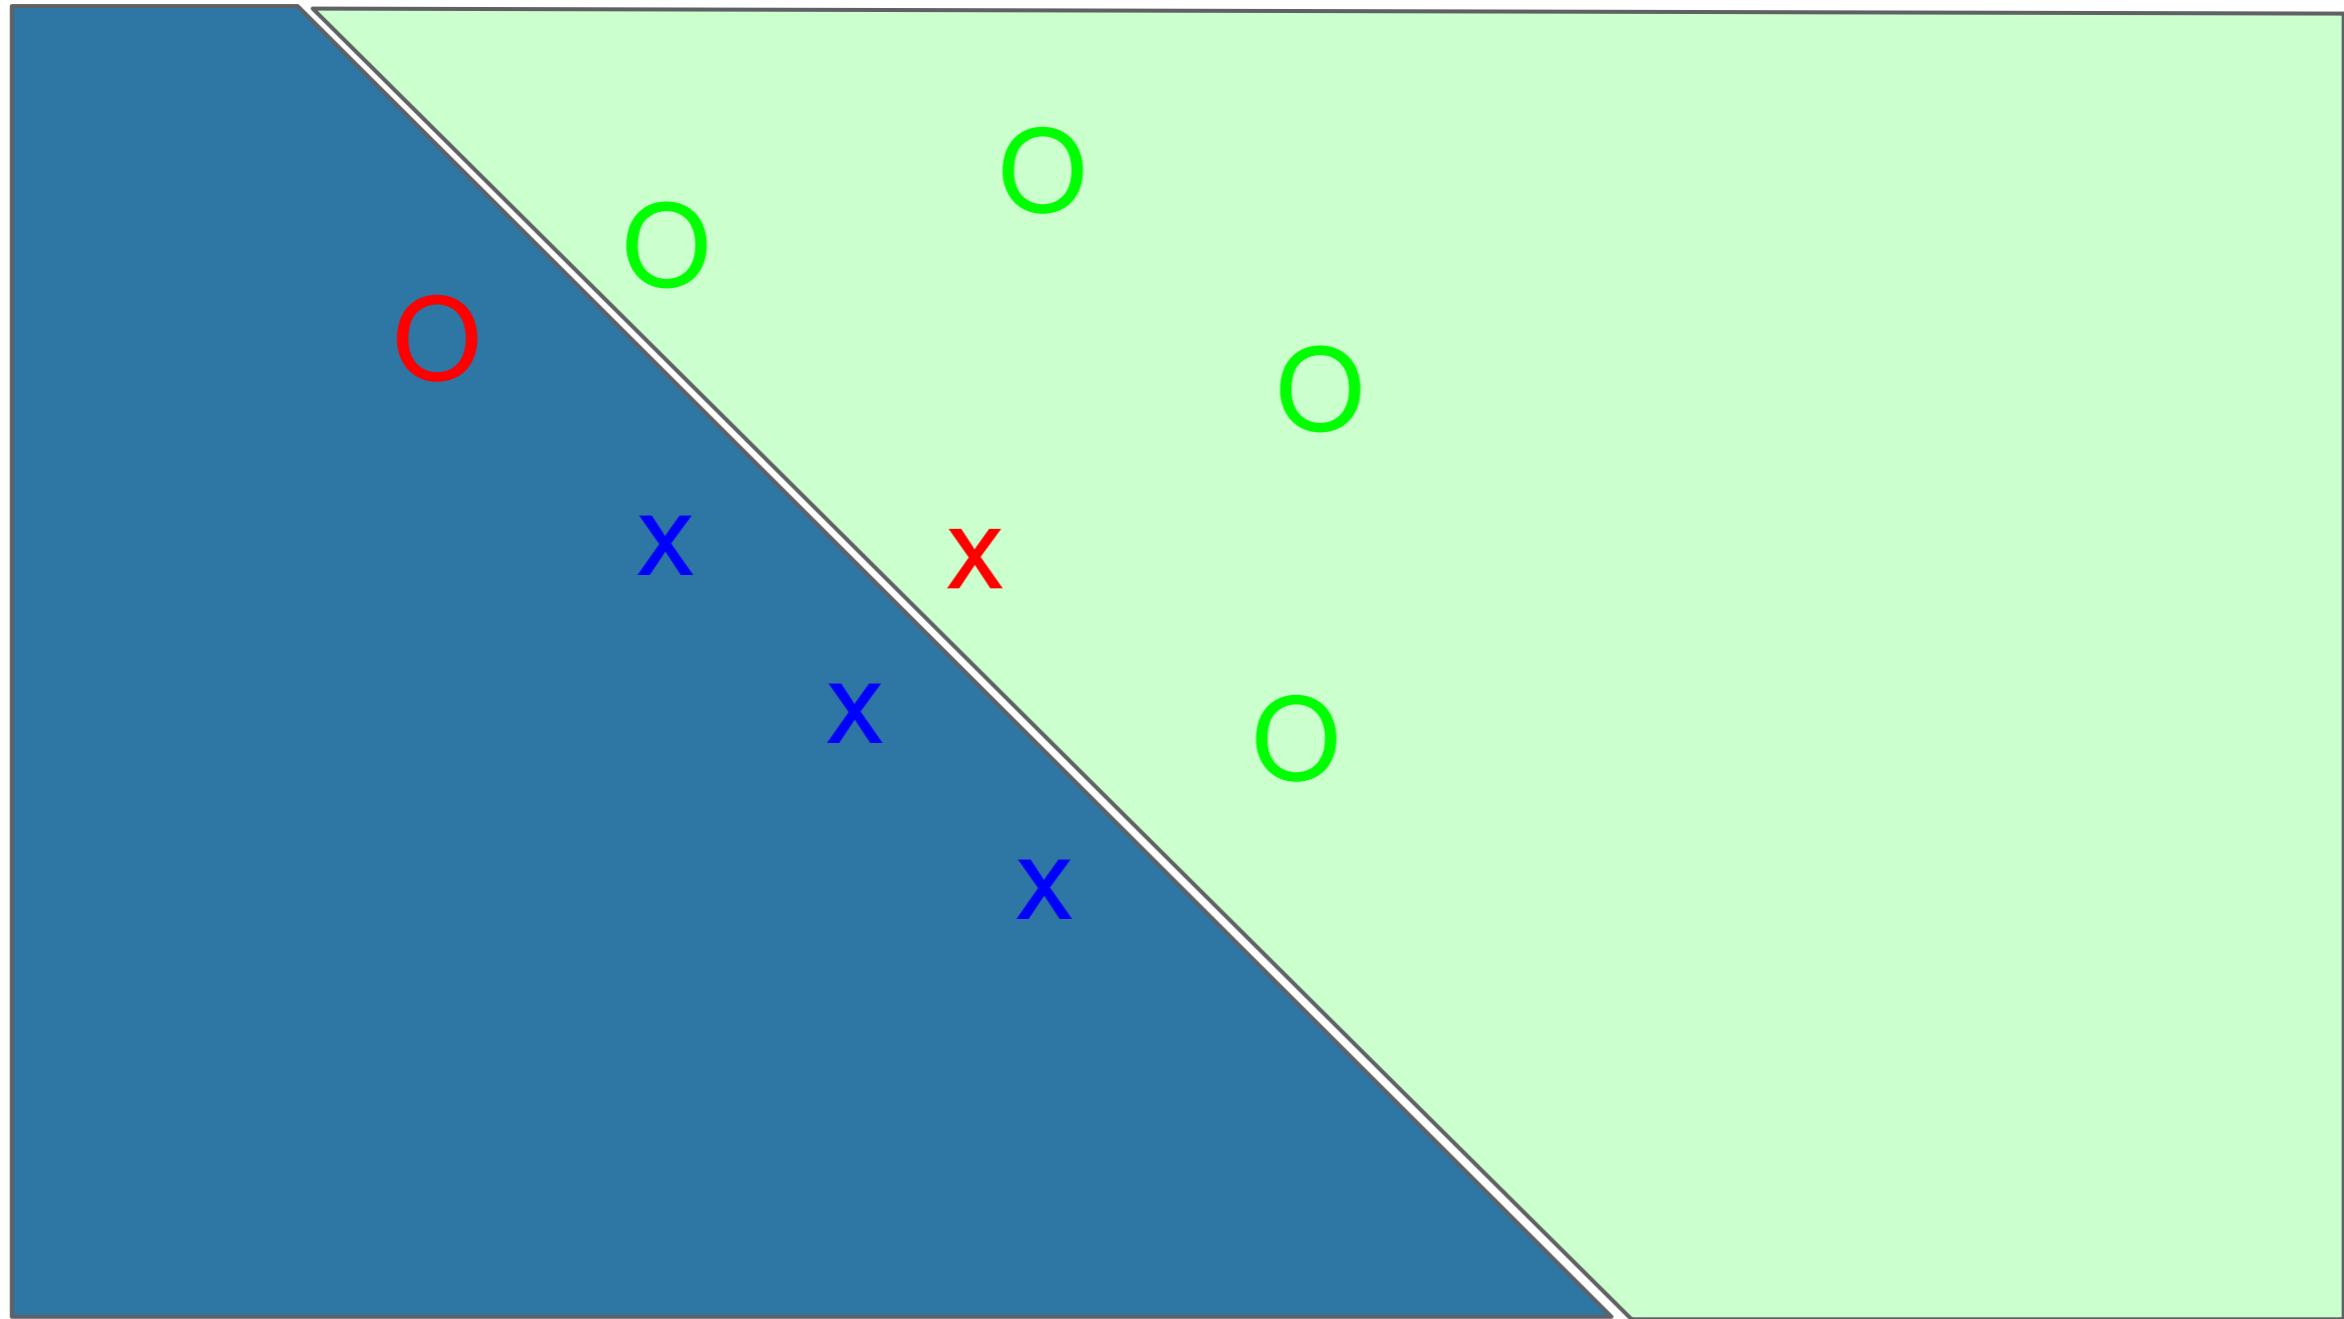
Not just for neural nets

- Linear models
 - Logistic regression
 - Softmax regression
 - SVMs
- Decision trees
- Nearest neighbors

Adversarial Examples from Overfitting

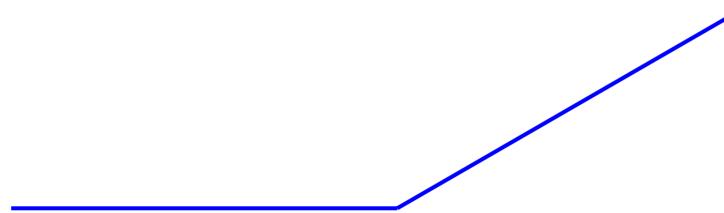


Adversarial Examples from Excessive Linearity

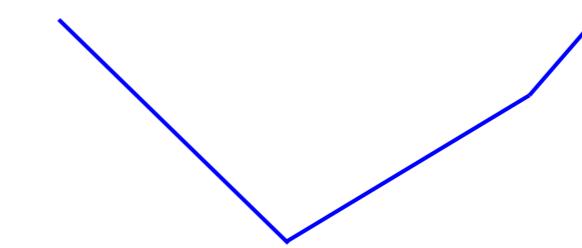


Modern deep nets are very piecewise linear

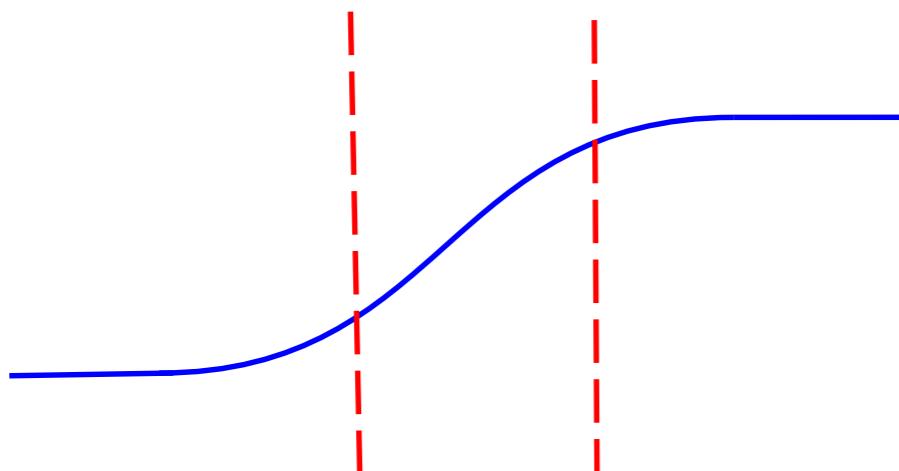
Rectified linear unit



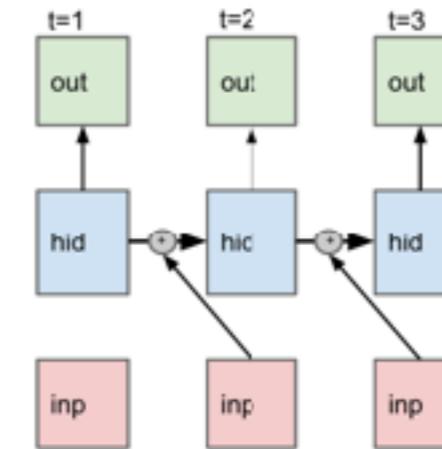
Maxout



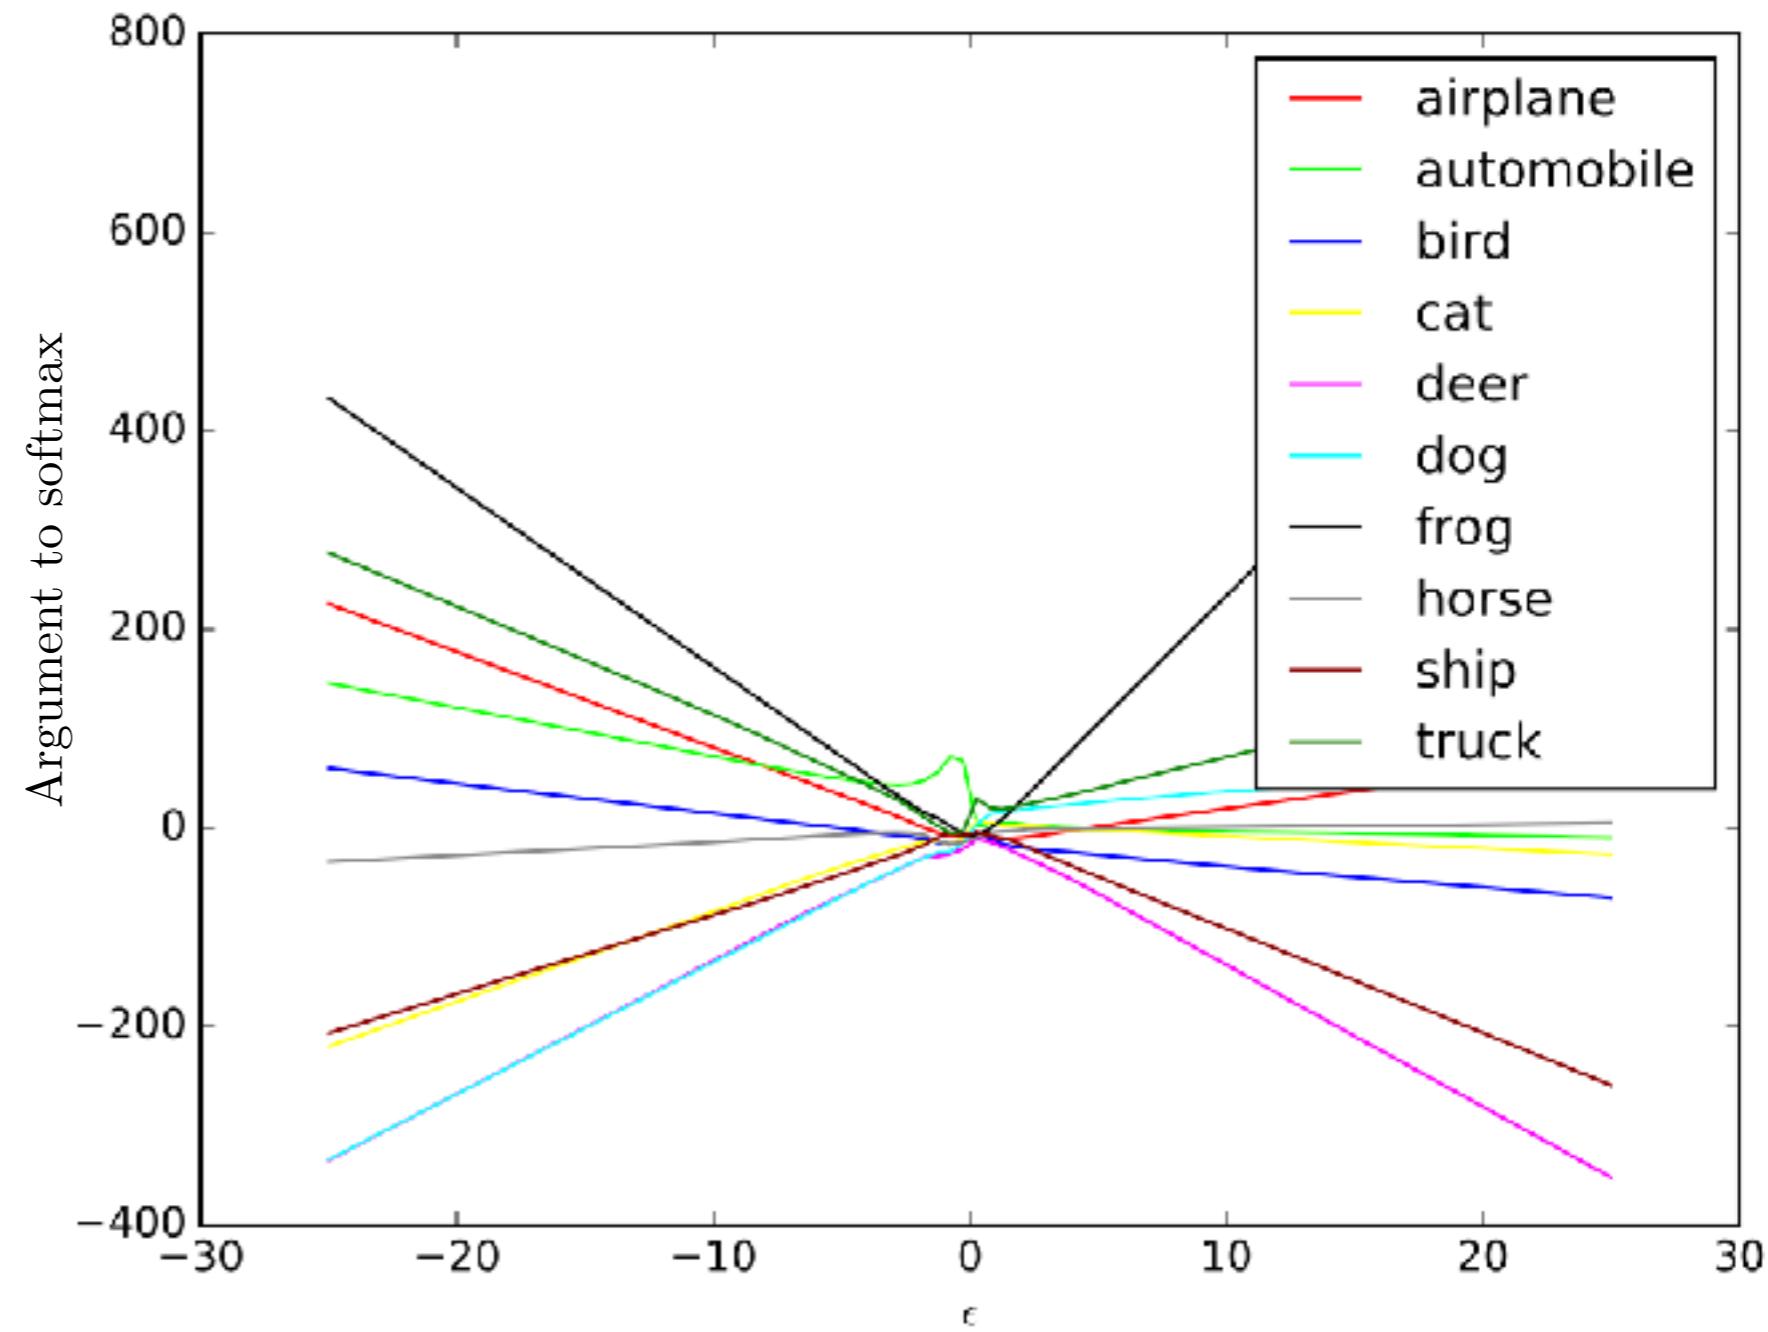
Carefully tuned sigmoid



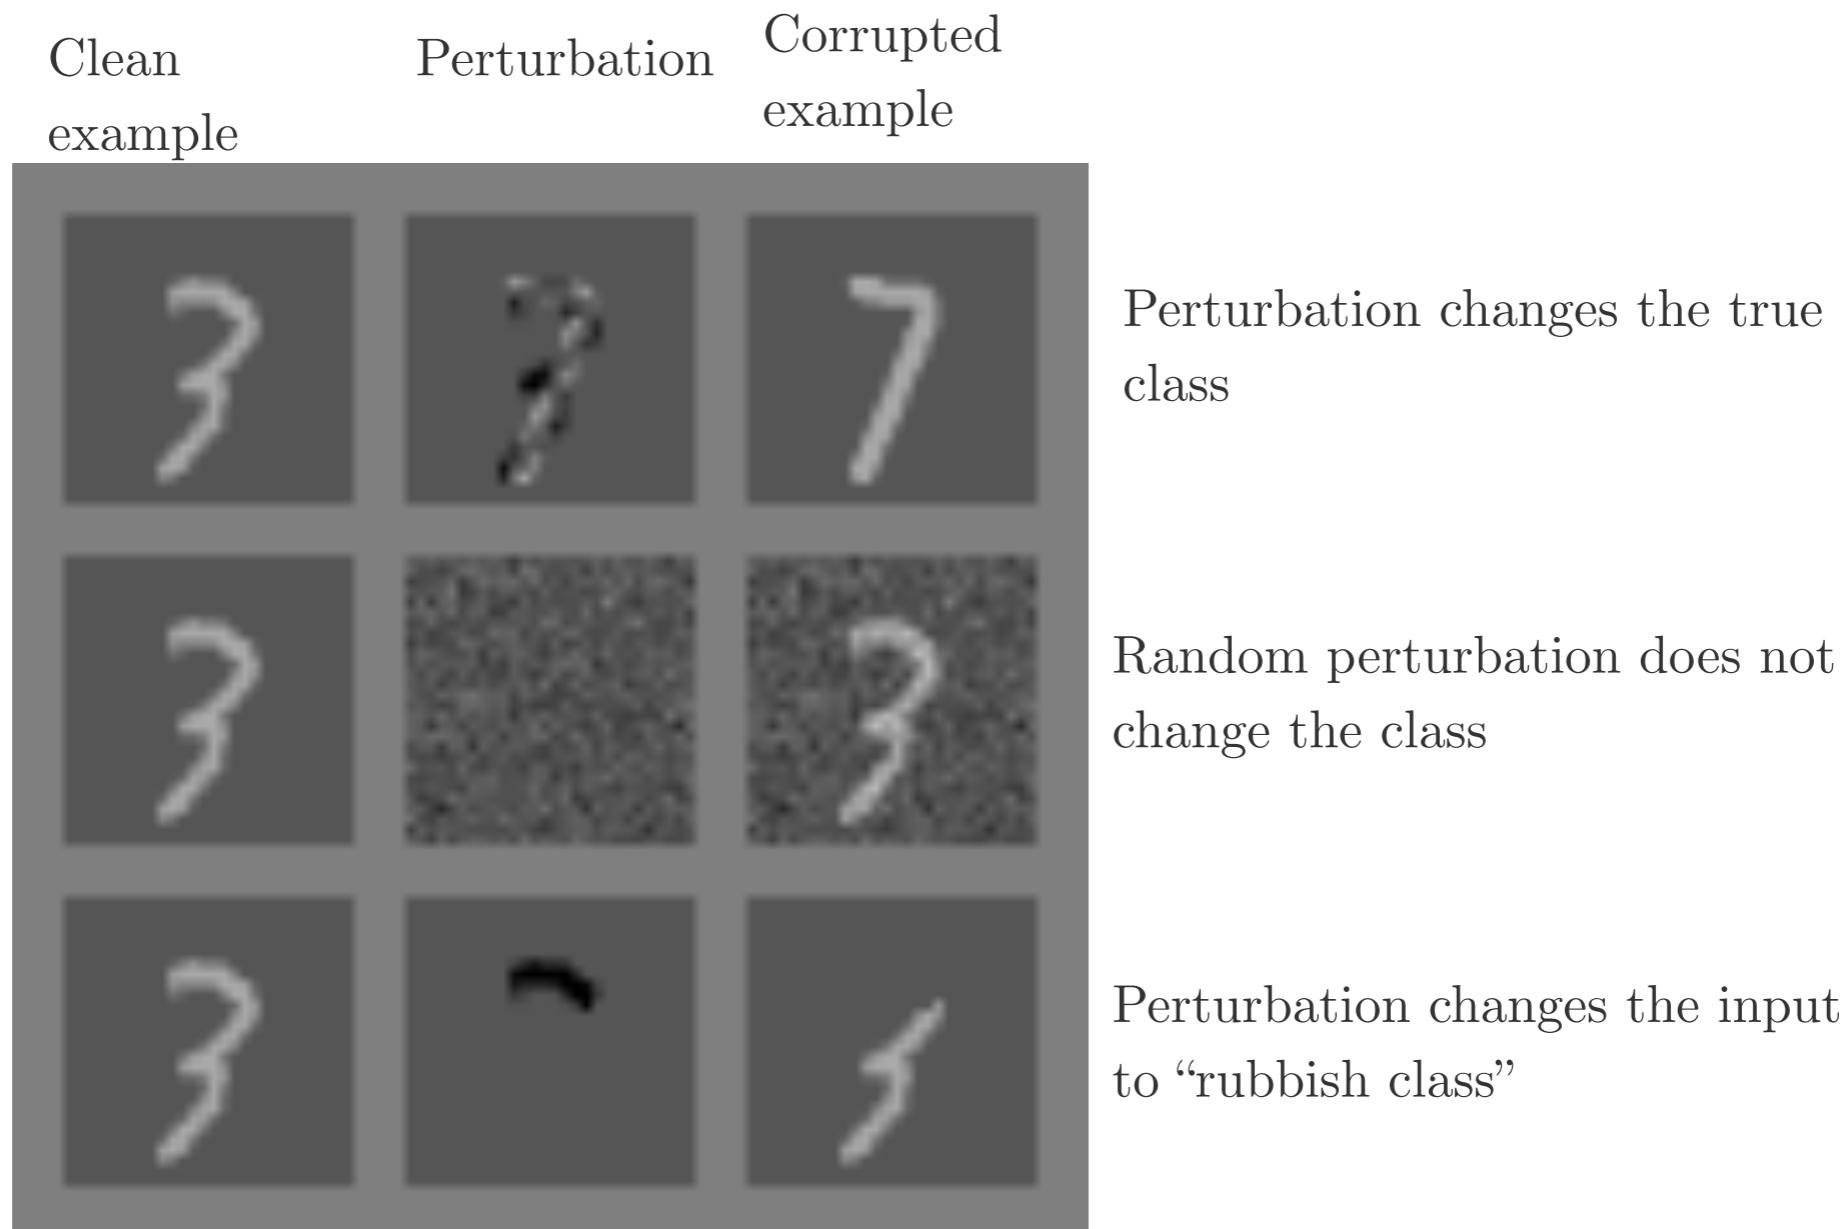
LSTM



Nearly Linear Responses in Practice



Small inter-class distances



All three perturbations have L2 norm 3.96

This is actually small. We typically use 7!

The Fast Gradient Sign Method

$$J(\tilde{\boldsymbol{x}}, \theta) \approx J(\boldsymbol{x}, \theta) + (\tilde{\boldsymbol{x}} - \boldsymbol{x})^\top \nabla_{\boldsymbol{x}} J(\boldsymbol{x}).$$

Maximize

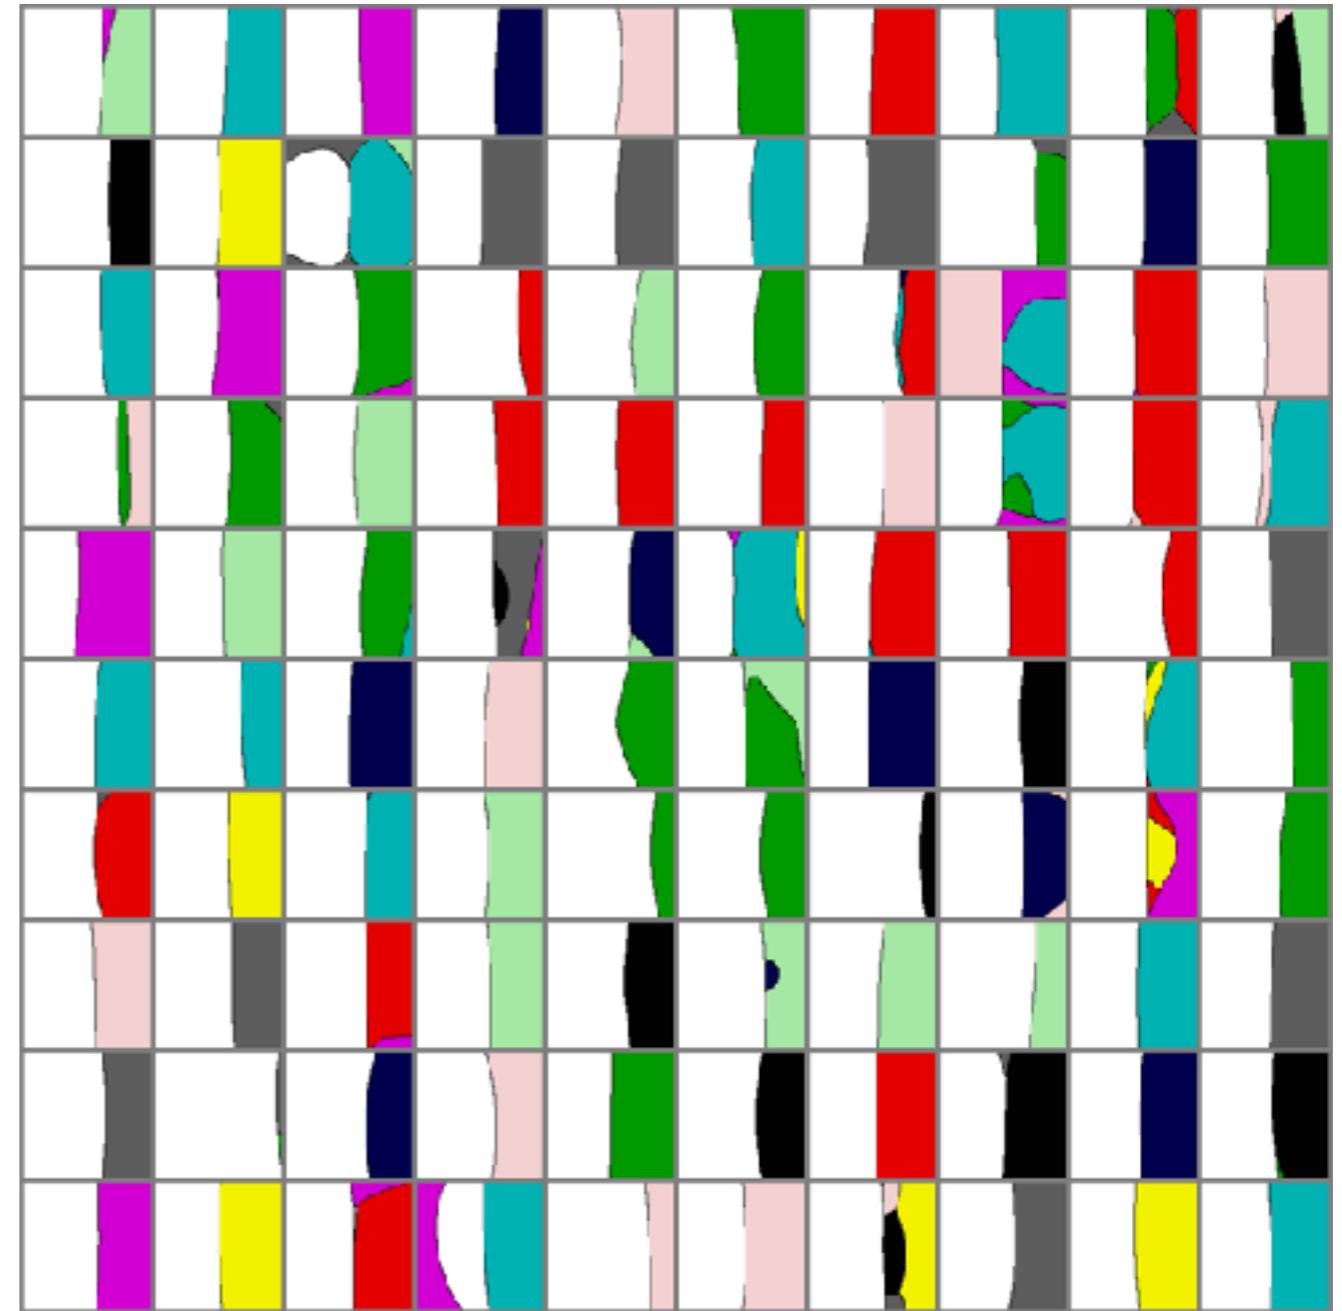
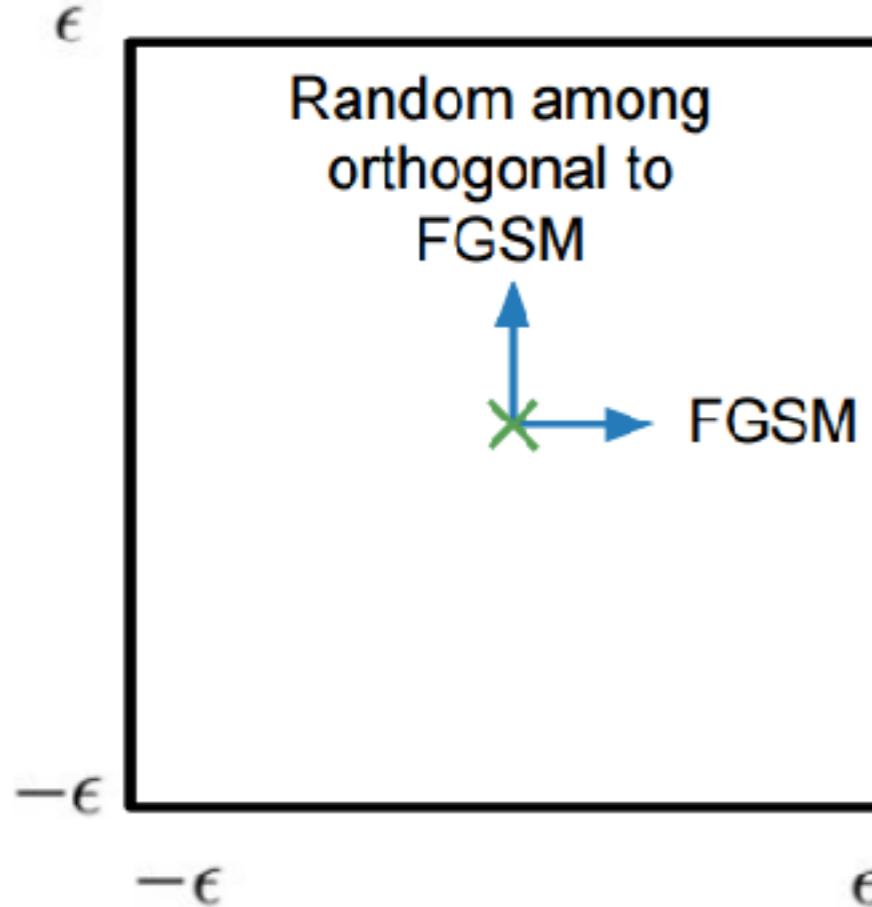
$$J(\boldsymbol{x}, \theta) + (\tilde{\boldsymbol{x}} - \boldsymbol{x})^\top \nabla_{\boldsymbol{x}} J(\boldsymbol{x})$$

subject to

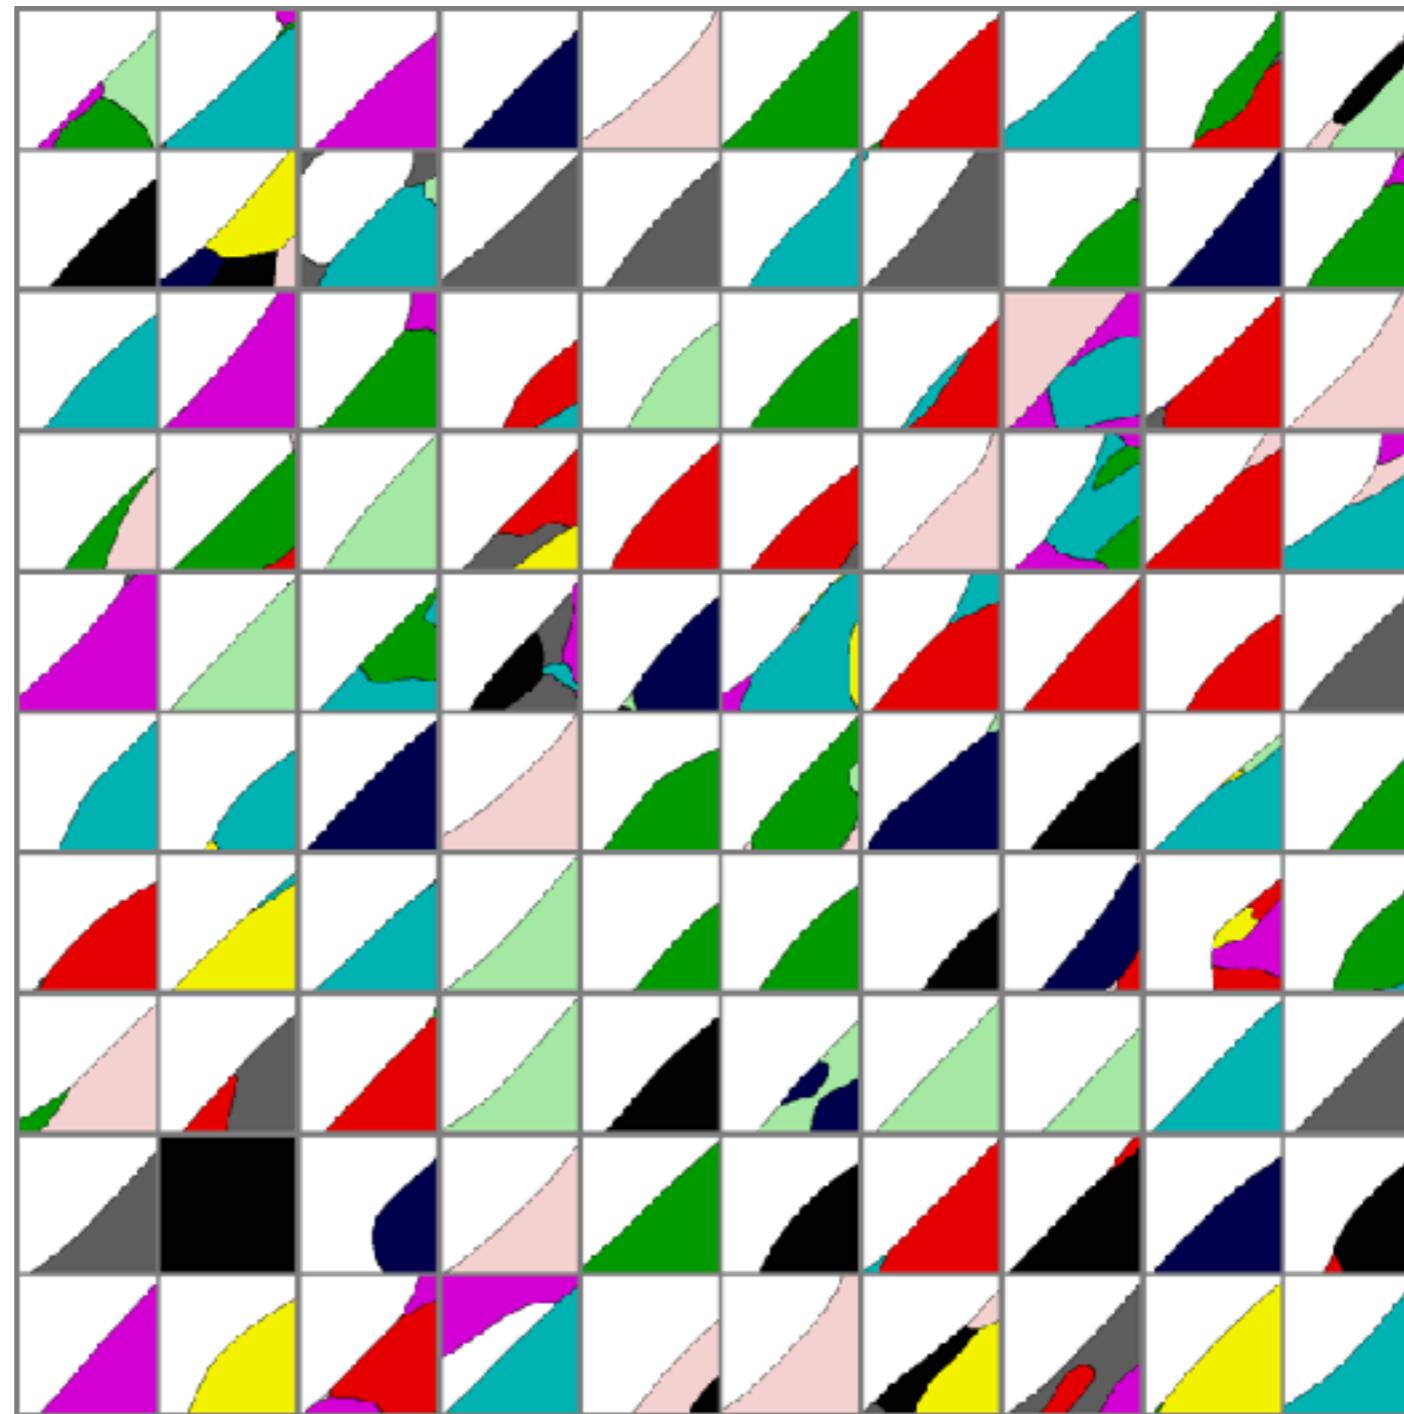
$$\|\tilde{\boldsymbol{x}} - \boldsymbol{x}\|_\infty \leq \epsilon$$

$$\Rightarrow \tilde{\boldsymbol{x}} = \boldsymbol{x} + \epsilon \text{sign}(\nabla_{\boldsymbol{x}} J(\boldsymbol{x})).$$

Maps of Adversarial and Random Cross-Sections

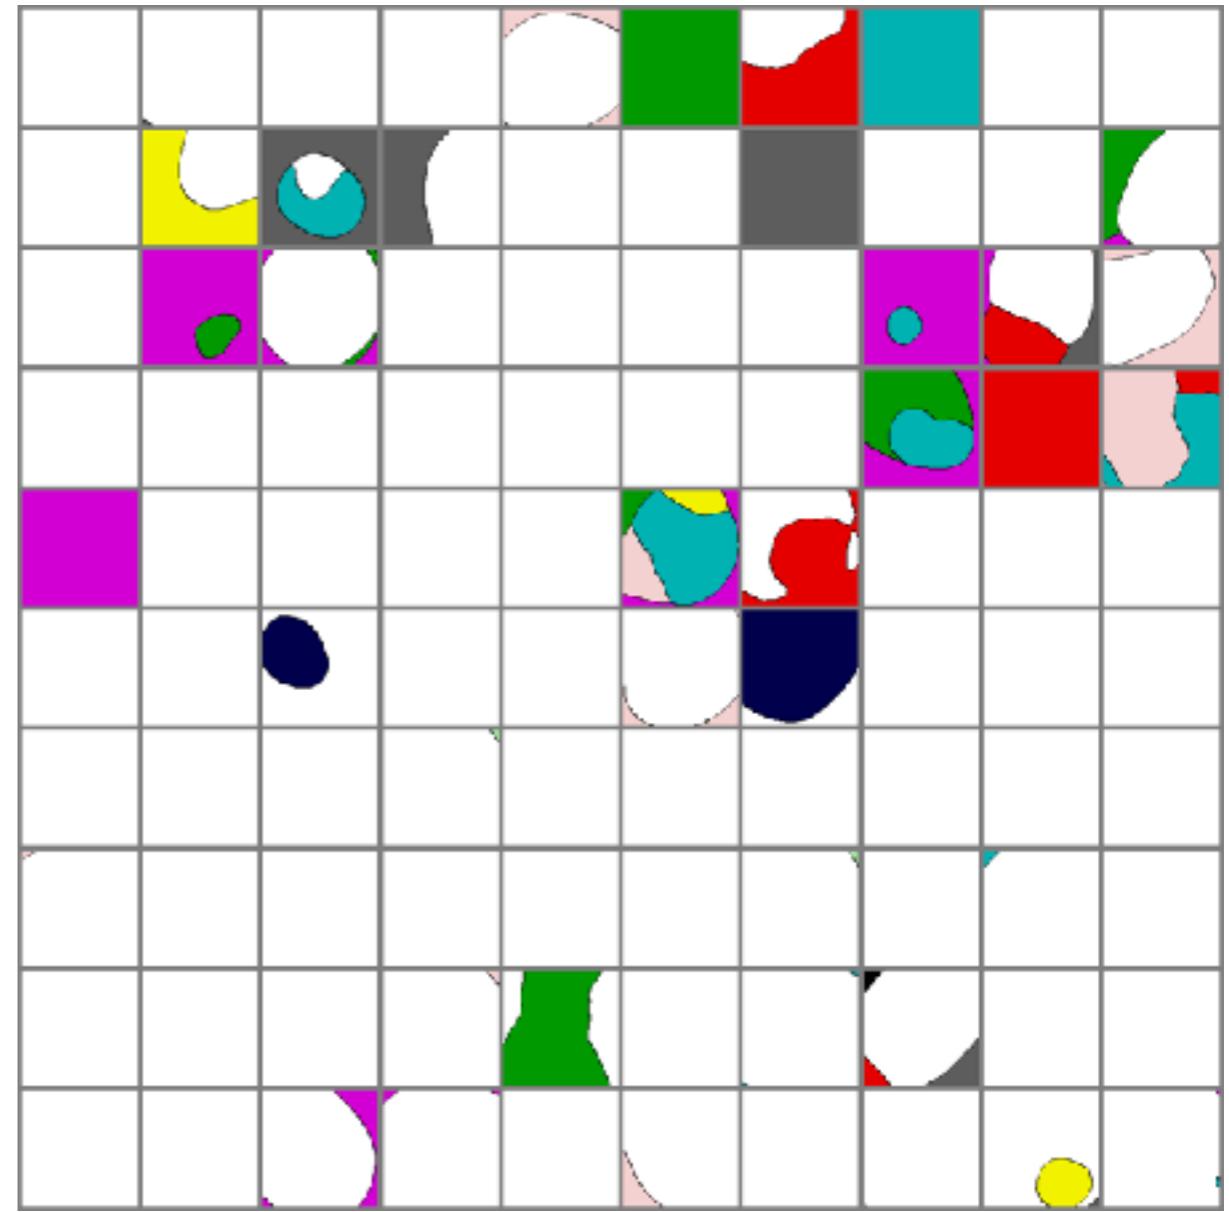
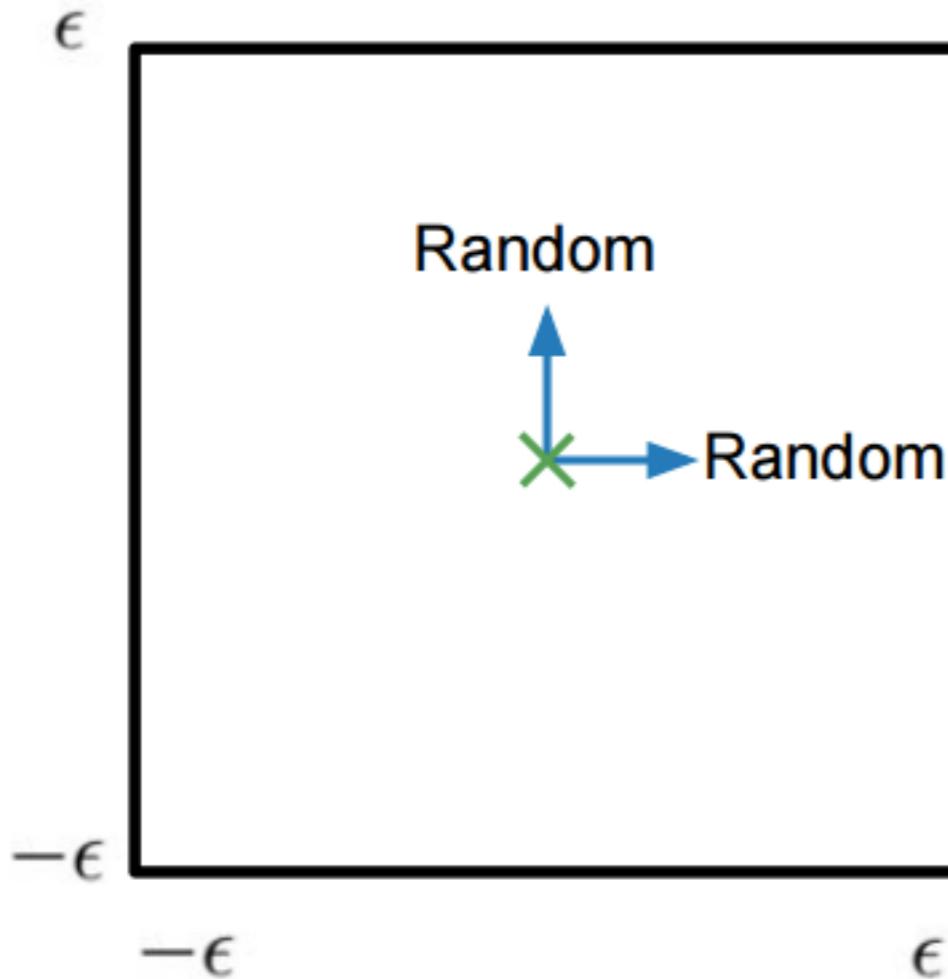


Maps of Adversarial Cross-Sections

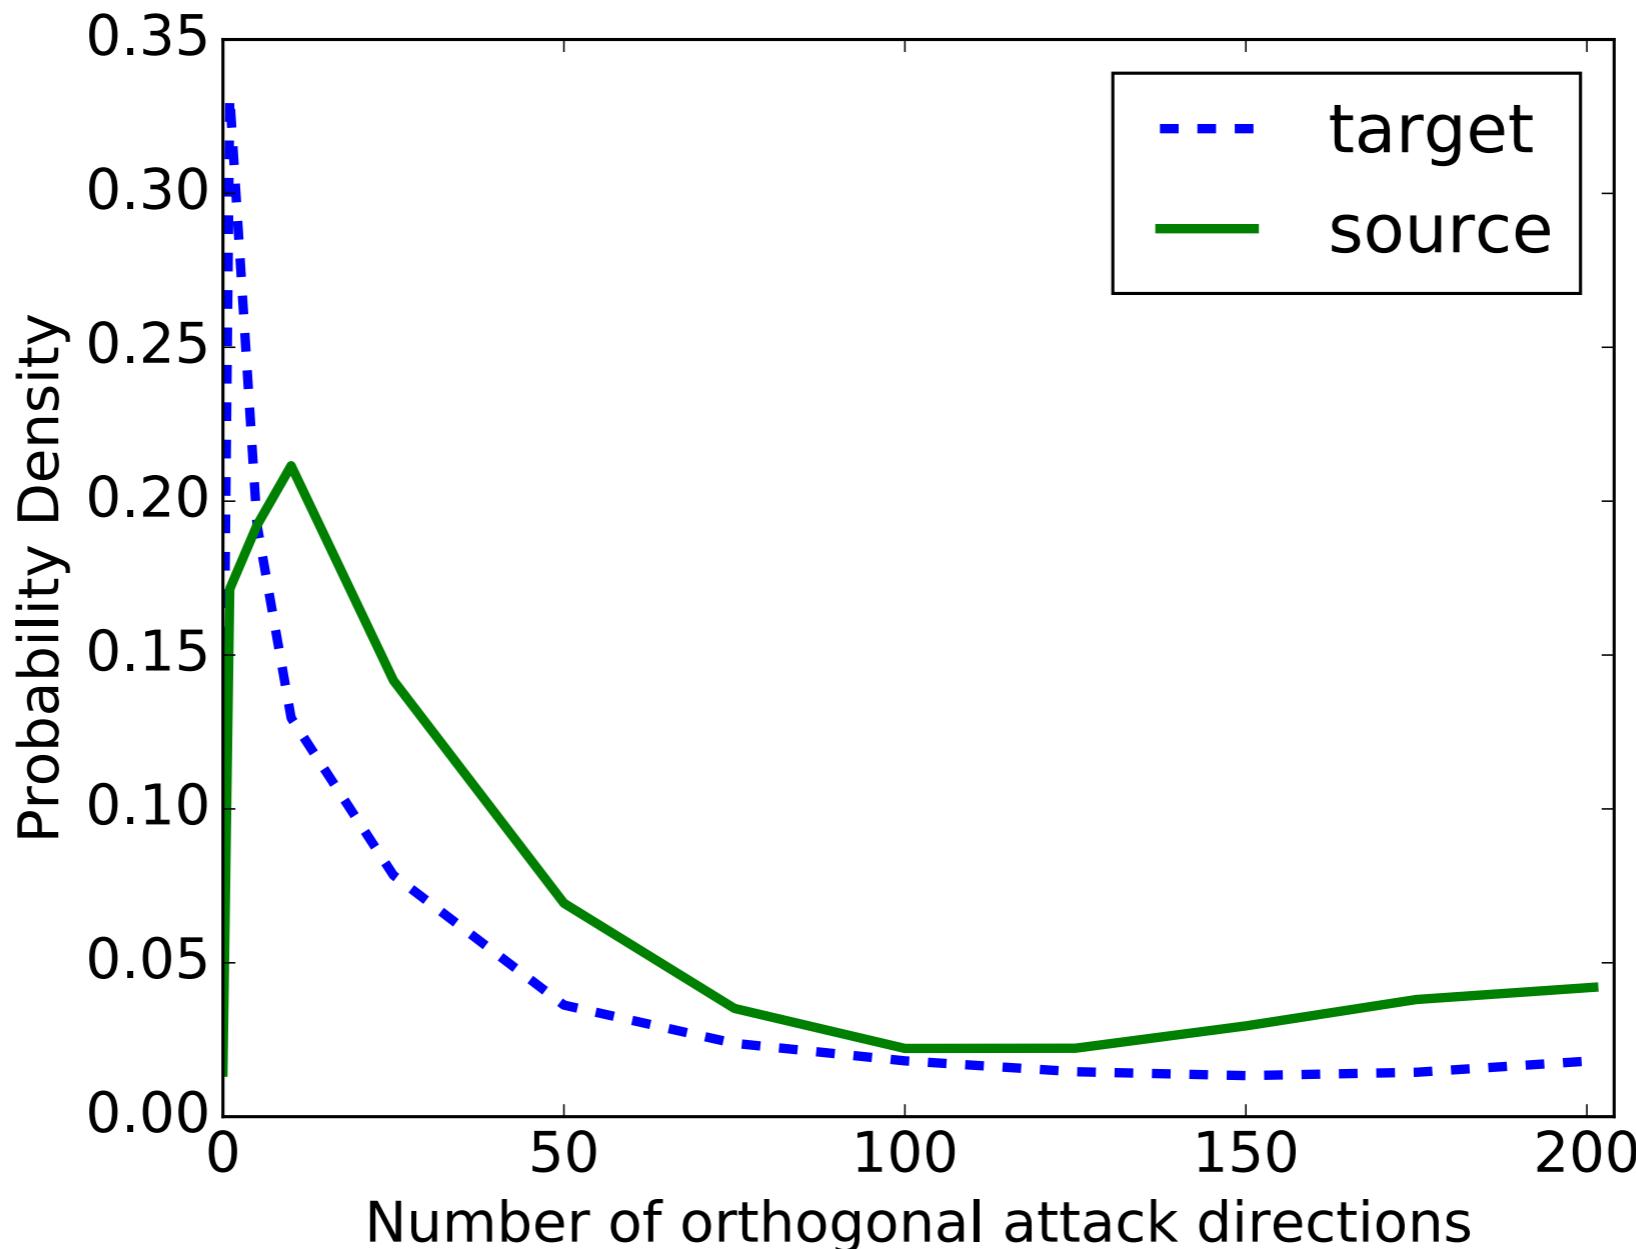


Maps of Random Cross-Sections

Adversarial examples
are not noise



Estimating the Subspace Dimensionality



(Tramèr et al, 2017)

(Goodfellow 2016)

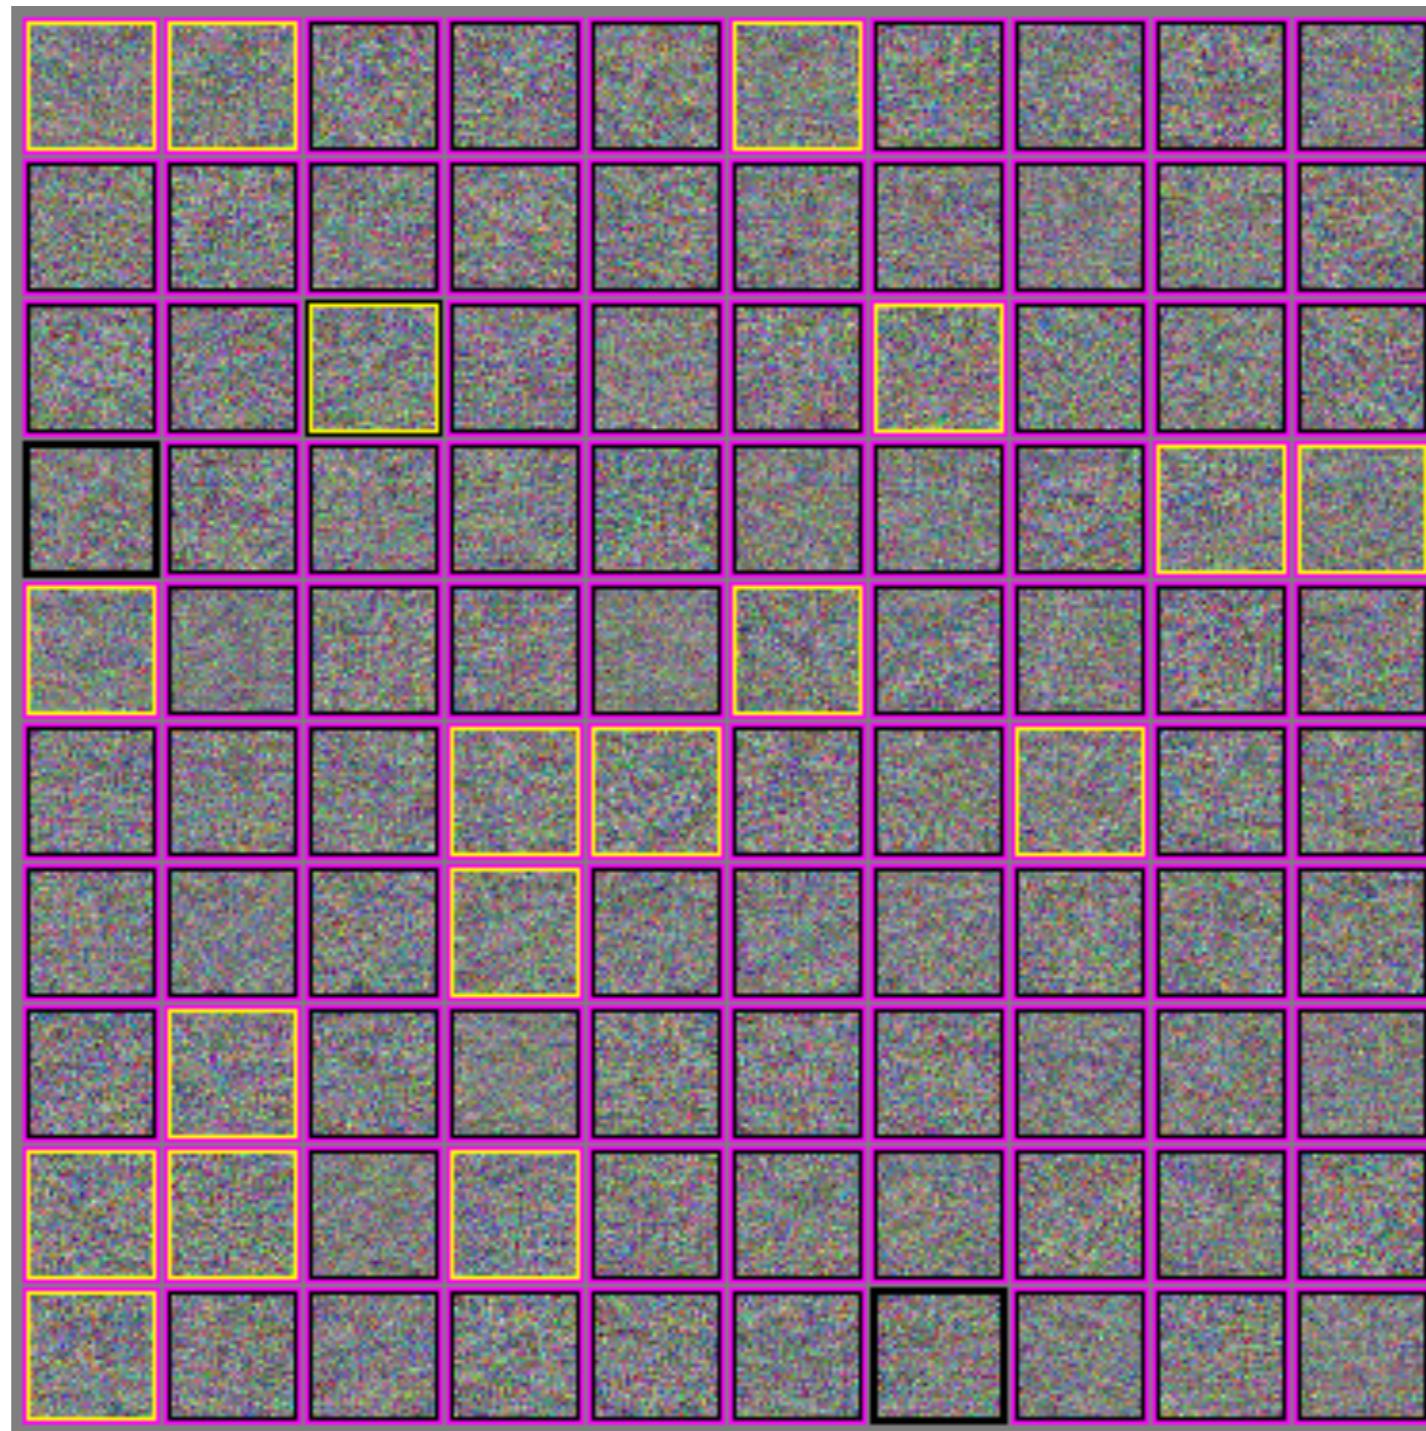
Clever Hans



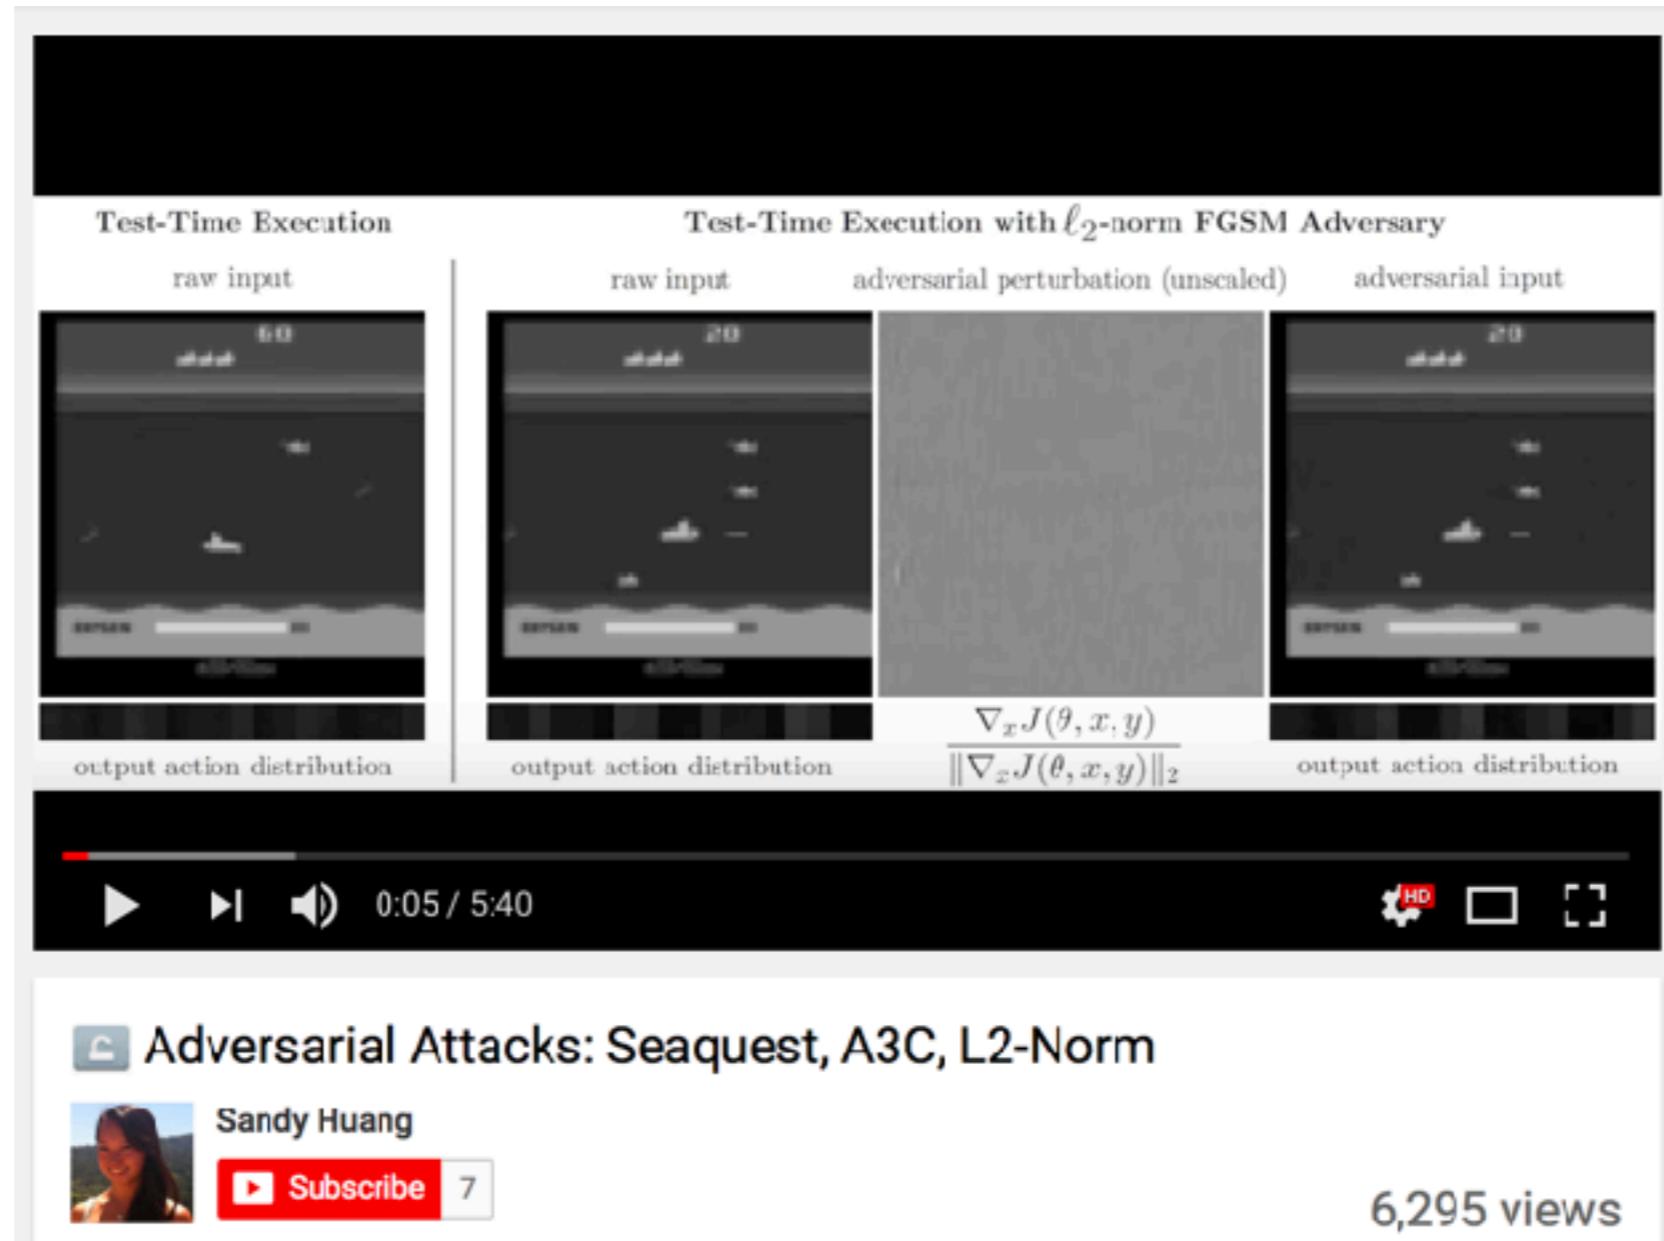
(“Clever Hans,
Clever
Algorithms,”
Bob Sturm)



Wrong almost everywhere

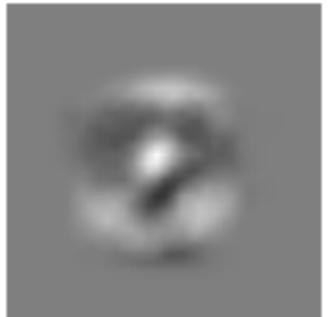
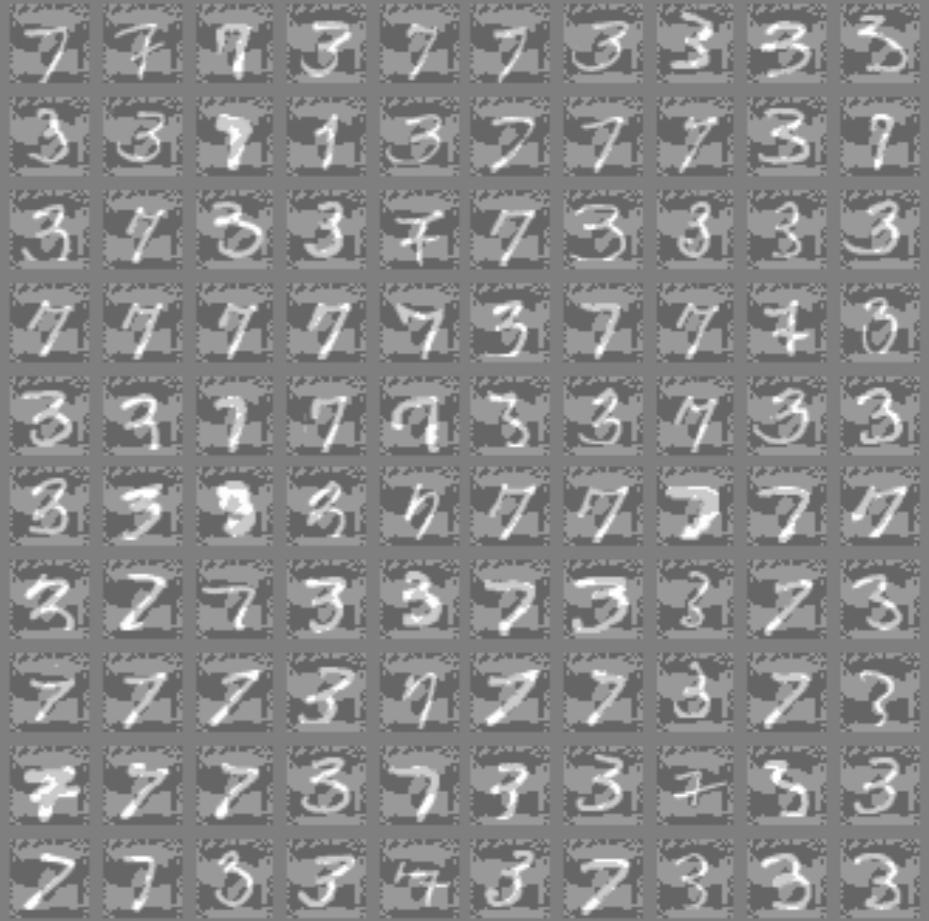
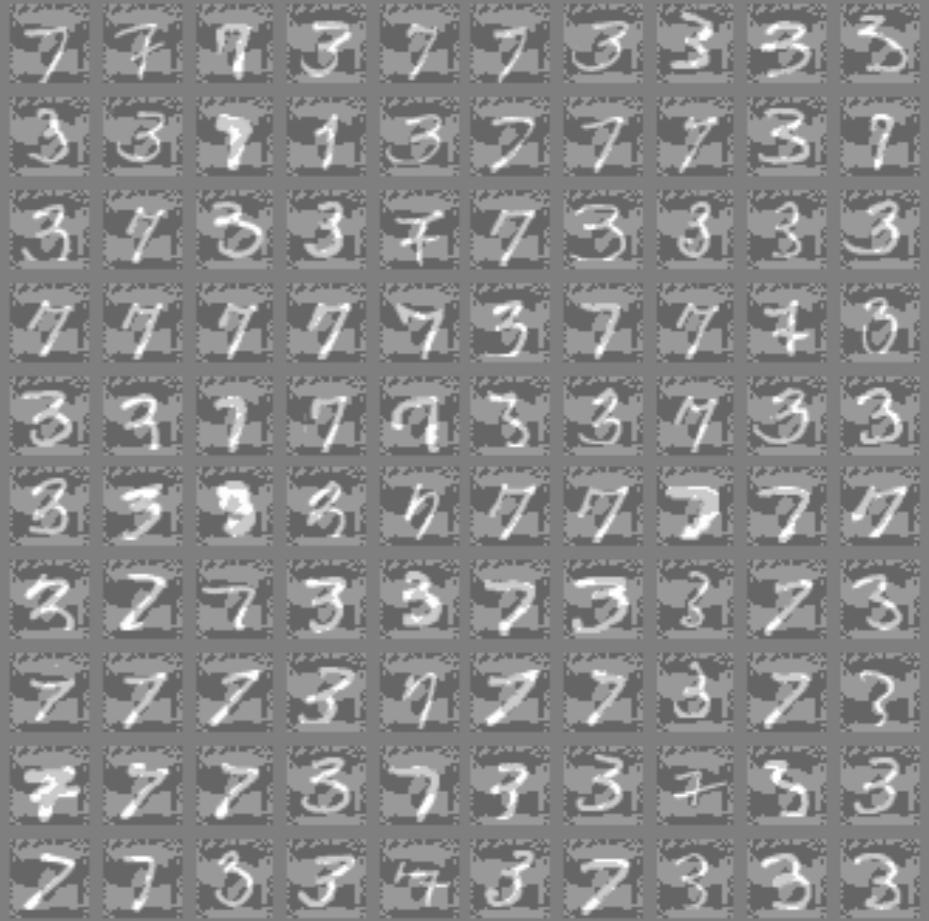


Adversarial Examples for RL



(Huang et al., 2017)

High-Dimensional Linear Models

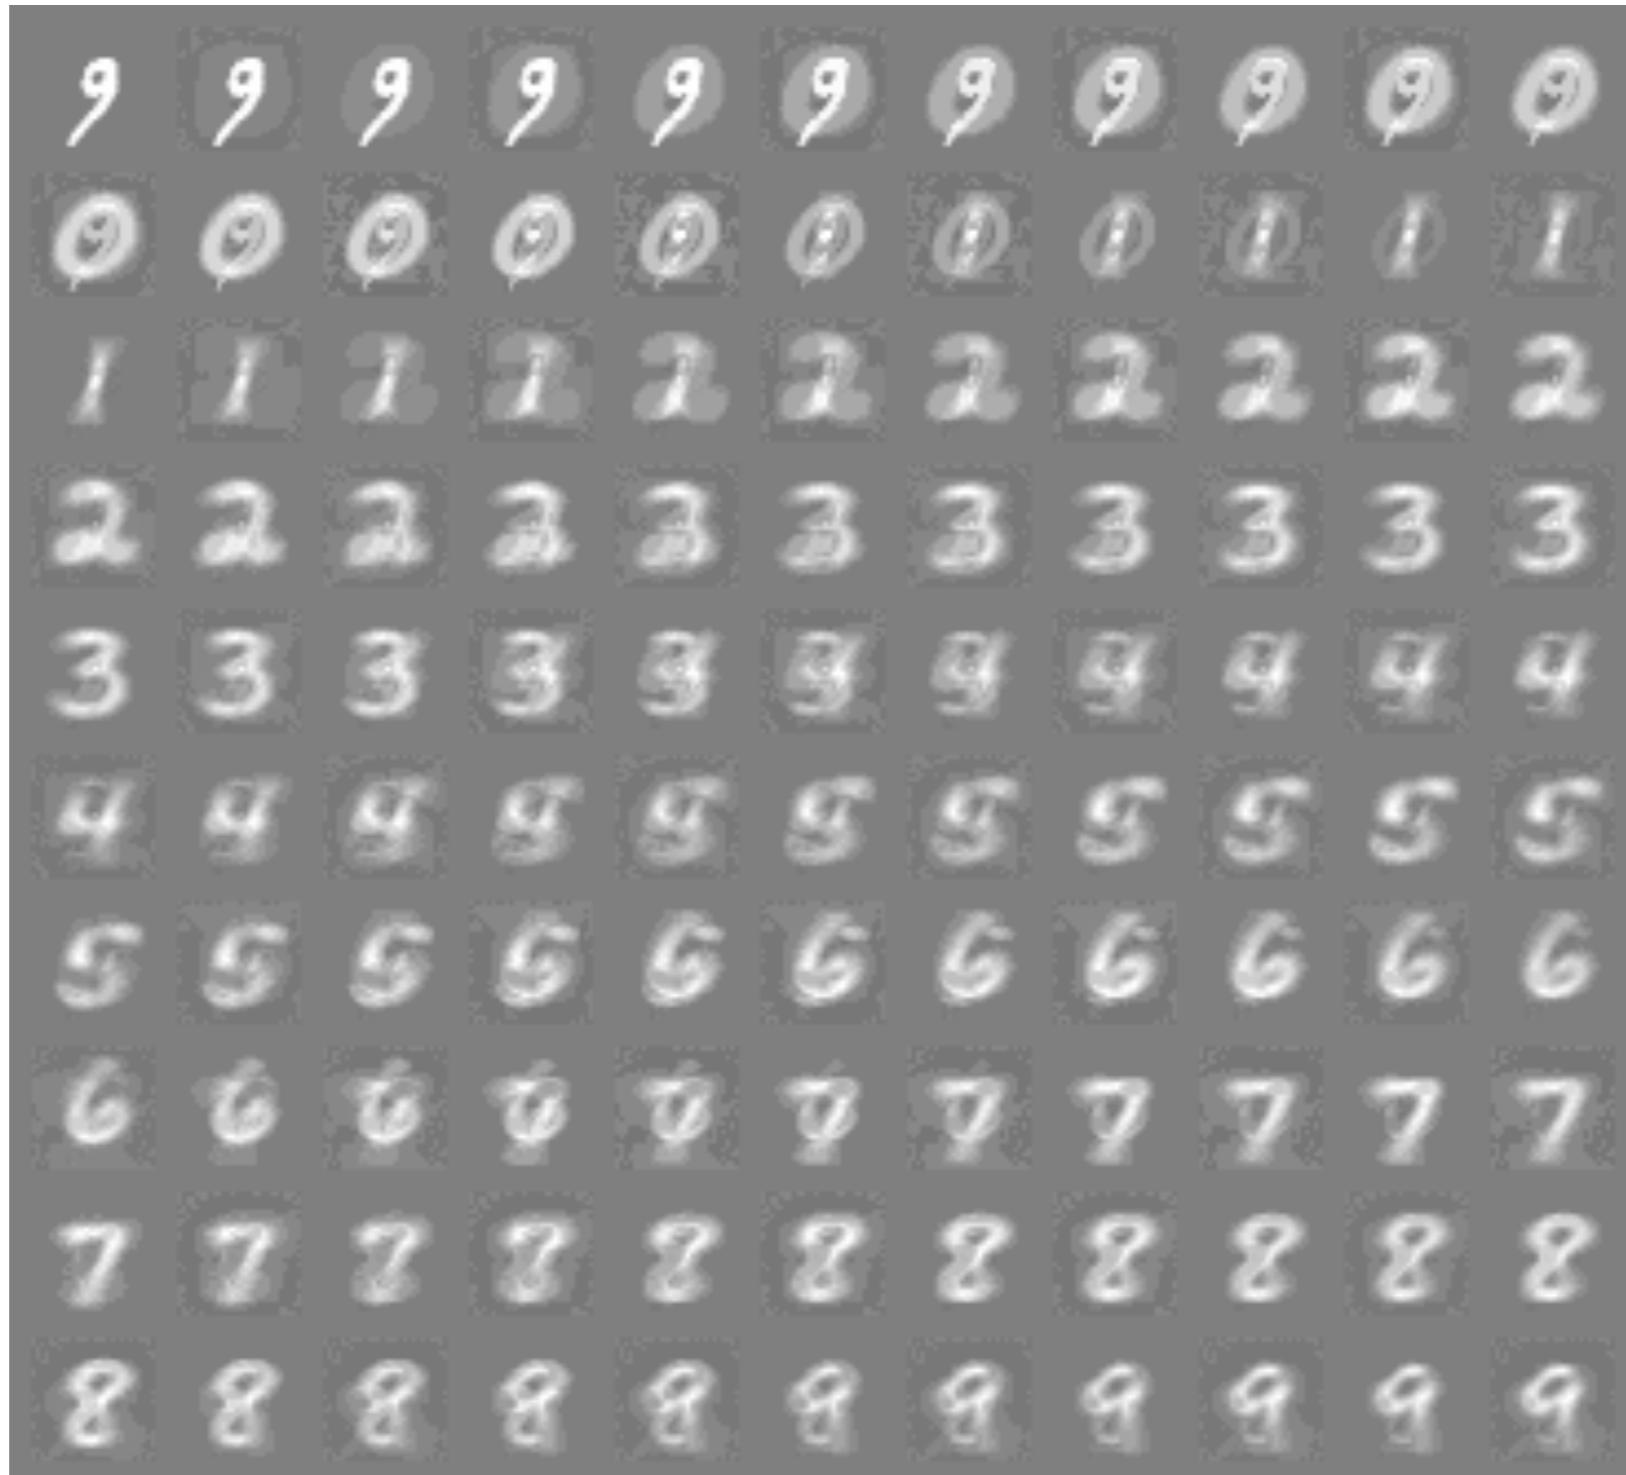
	Clean examples	Adversarial
Weights		
Signs of weights		

Linear Models of ImageNet

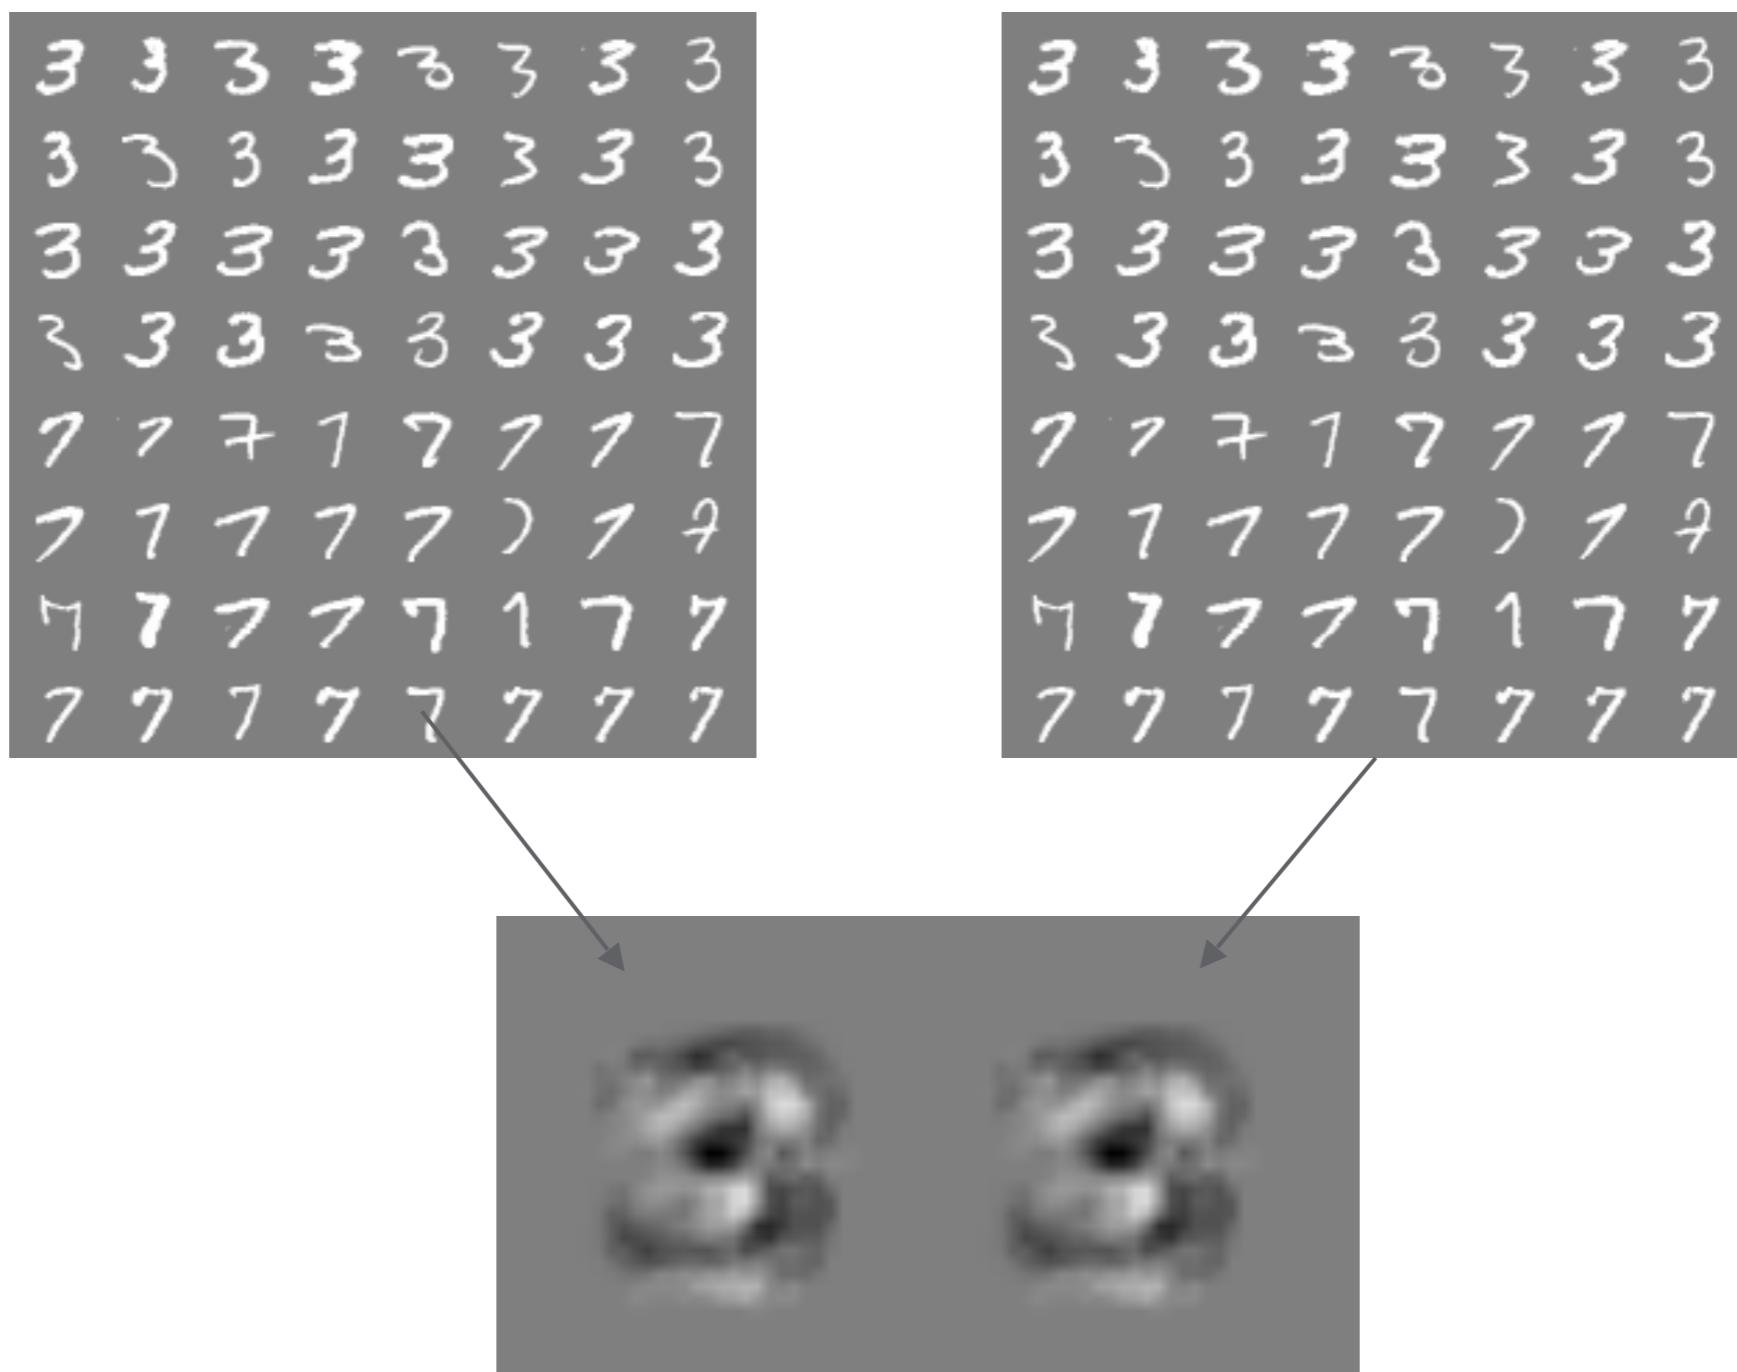


(Andrej Karpathy, “Breaking Linear Classifiers on ImageNet”)

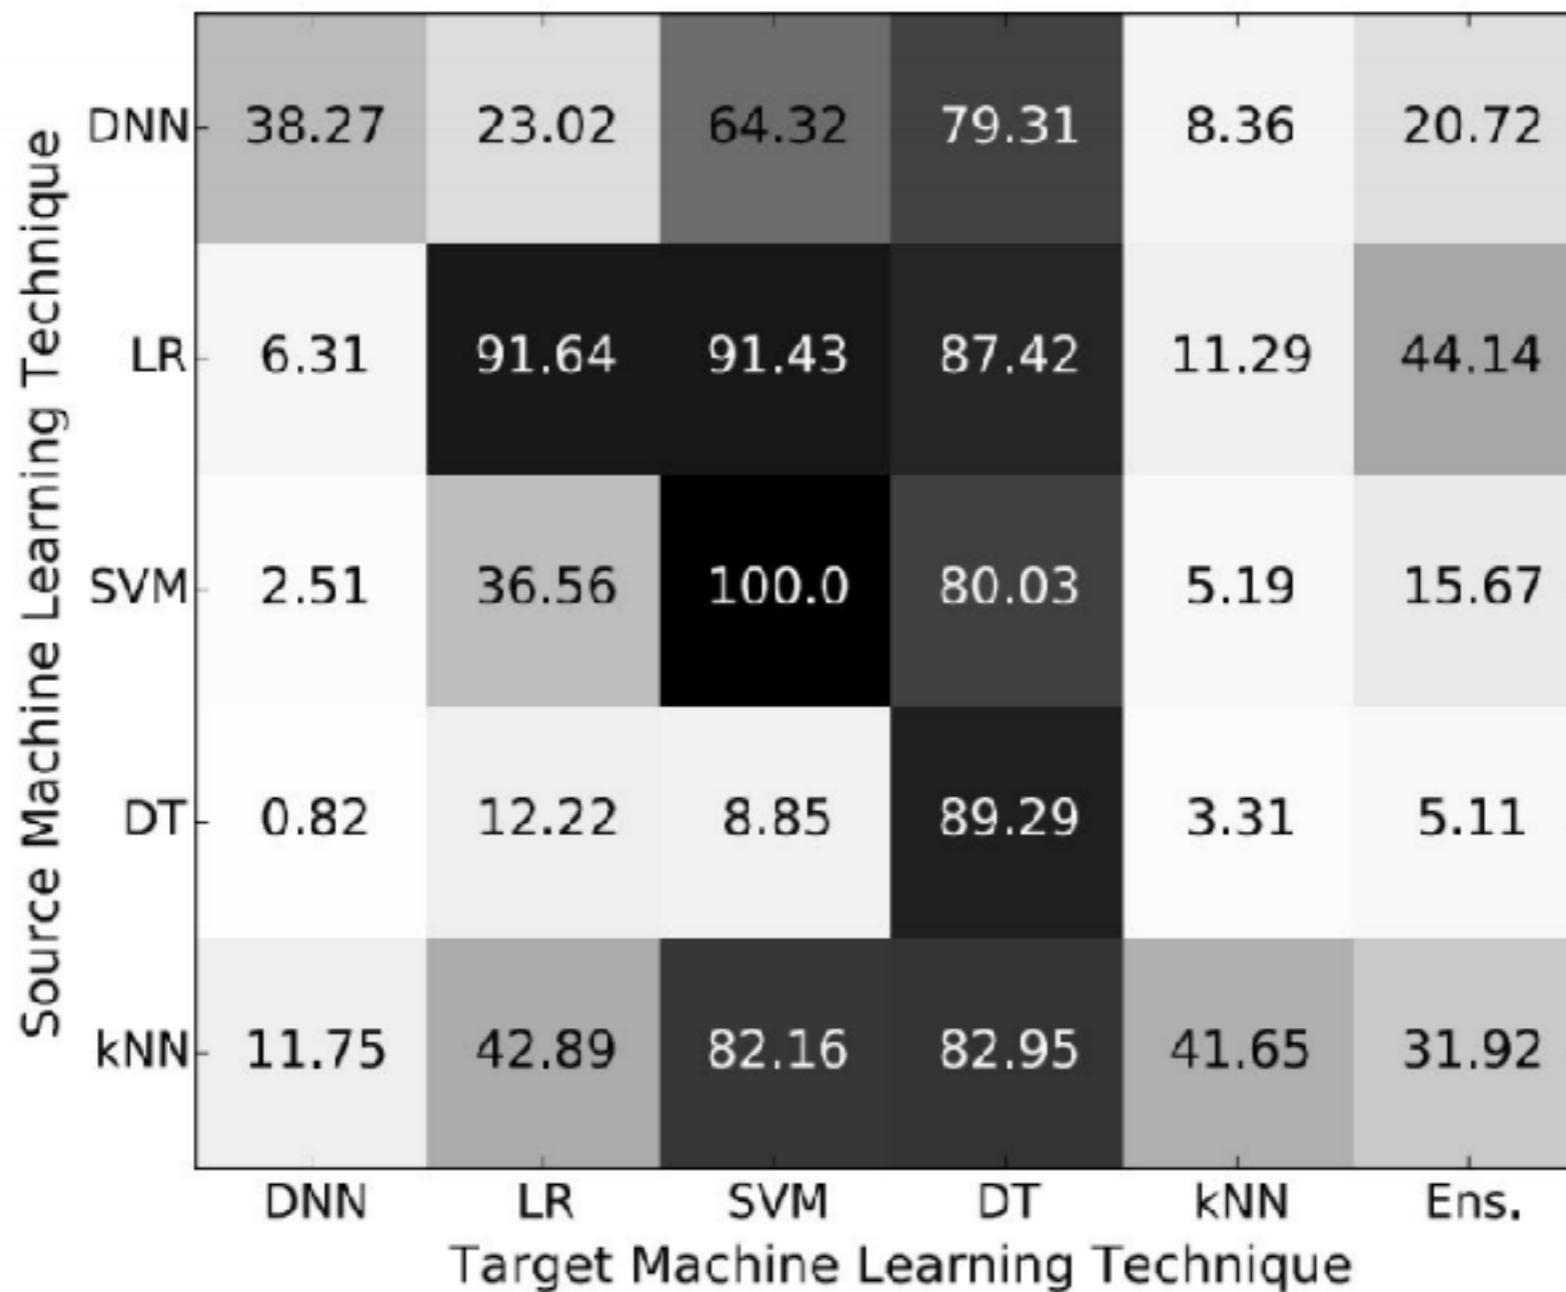
RBFs behave more intuitively



Cross-model, cross-dataset generalization



Cross-technique transferability



(Papernot 2016)

Transferability Attack

Target model with unknown weights, machine learning algorithm, training set; maybe non-differentiable

Train your own model

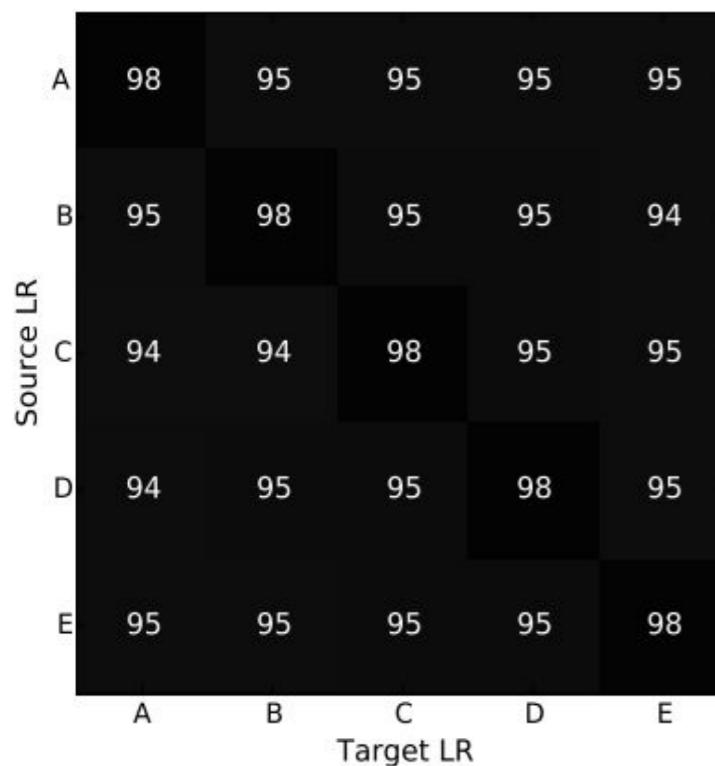
Substitute model mimicking target model with known, differentiable function

Deploy adversarial examples against the target; transferability property results in them succeeding

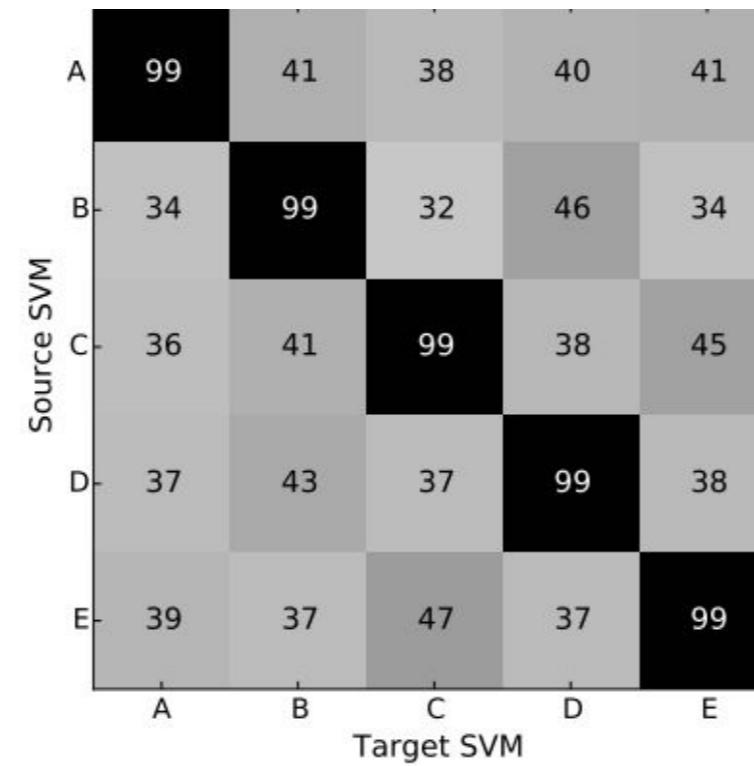
Adversarial examples

Adversarial crafting against substitute

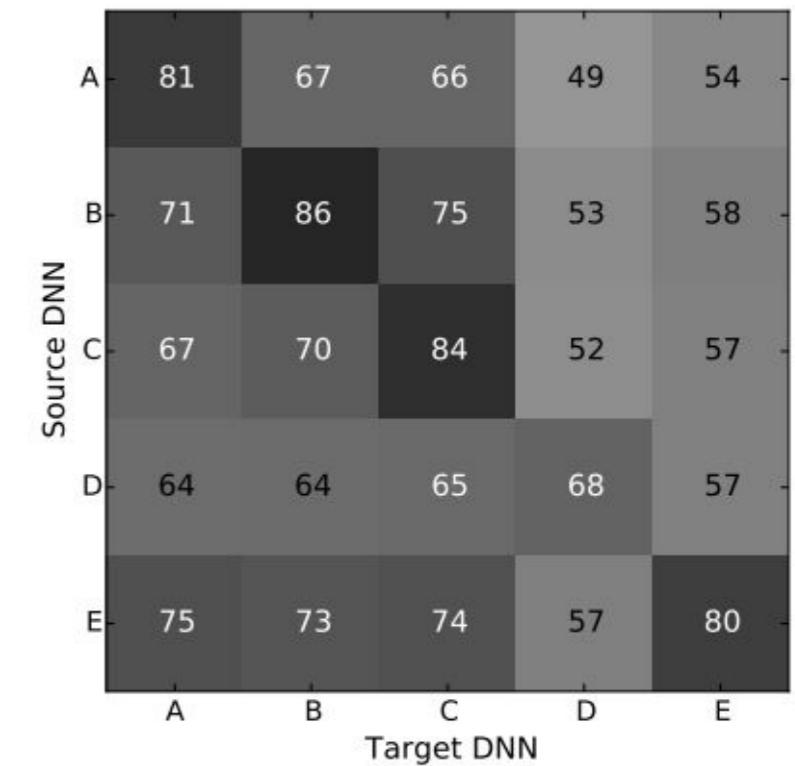
Cross-Training Data Transferability



Strong



Weak



Intermediate

(Papernot 2016)

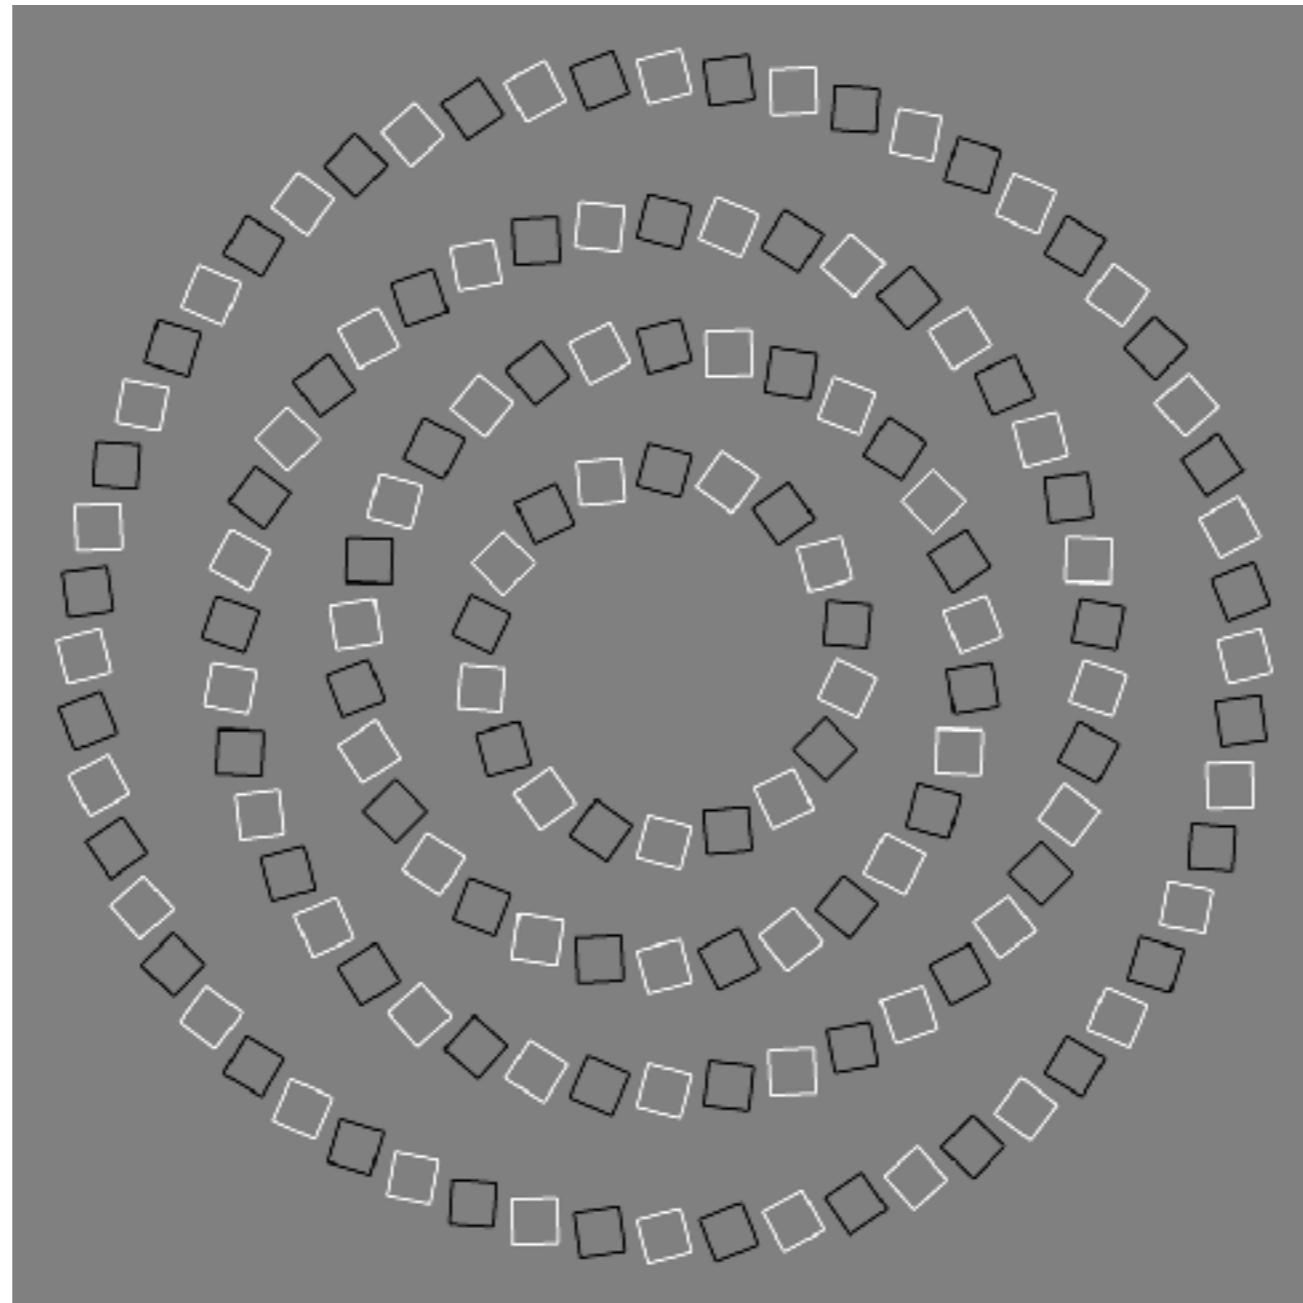
Enhancing Transfer With Ensembles

	RMSD	ResNet-152	ResNet-101	ResNet-50	VGG-16	GoogLeNet
-ResNet-152	17.17	0%	0%	0%	0%	0%
-ResNet-101	17.25	0%	1%	0%	0%	0%
-ResNet-50	17.25	0%	0%	2%	0%	0%
-VGG-16	17.80	0%	0%	0%	6%	0%
-GoogLeNet	17.41	0%	0%	0%	0%	5%

Table 4: Accuracy of non-targeted adversarial images generated using the optimization-based approach. The first column indicates the average RMSD of the generated adversarial images. Cell (i, j) corresponds to the accuracy of the attack generated using four models except model i (row) when evaluated over model j (column). In each row, the minus sign “-” indicates that the model of the row is not used when generating the attacks. Results of top-5 accuracy can be found in the appendix (Table 14).

(Liu et al, 2016)

Adversarial Examples in the Human Brain



These are
concentric
circles,
not
intertwined
spirals.

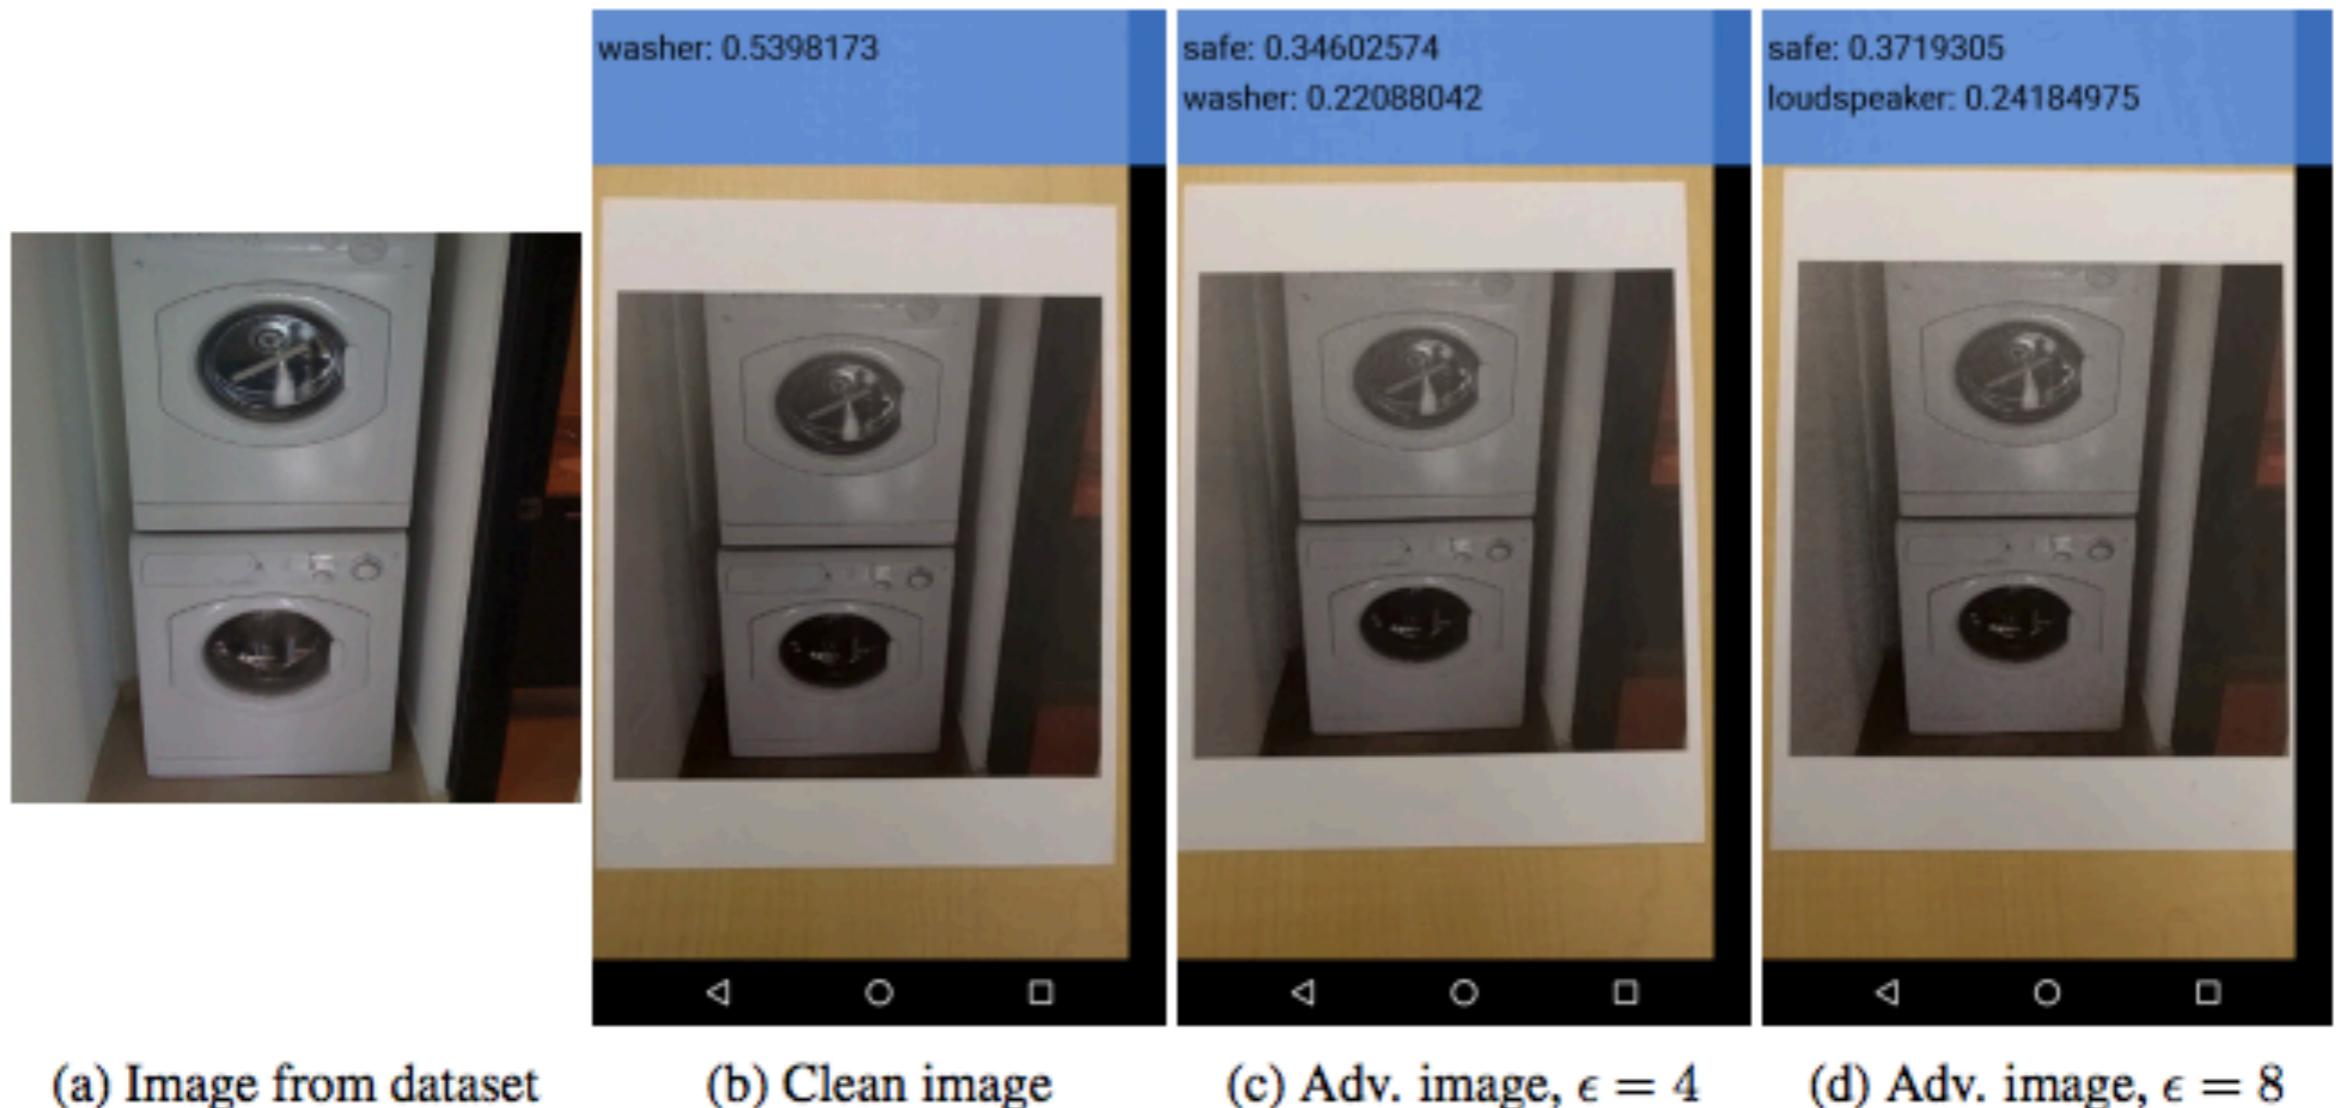
(Pinna and Gregory, 2002)

(Goodfellow 2016)

Practical Attacks

- Fool real classifiers trained by remotely hosted API
(MetaMind, Amazon, Google)
- Fool malware detector networks
- Display adversarial examples in the physical world
and fool machine learning systems that perceive
them through a camera

Adversarial Examples in the Physical World



(a) Image from dataset

(b) Clean image

(c) Adv. image, $\epsilon = 4$

(d) Adv. image, $\epsilon = 8$

(Kurakin et al, 2016)

(Goodfellow 2016)

Failed defenses

Generative
pretraining

Adding noise
at test time

Confidence-reducing
perturbation at test time

Weight decay

Various
non-linear units

Double backprop
Dropout

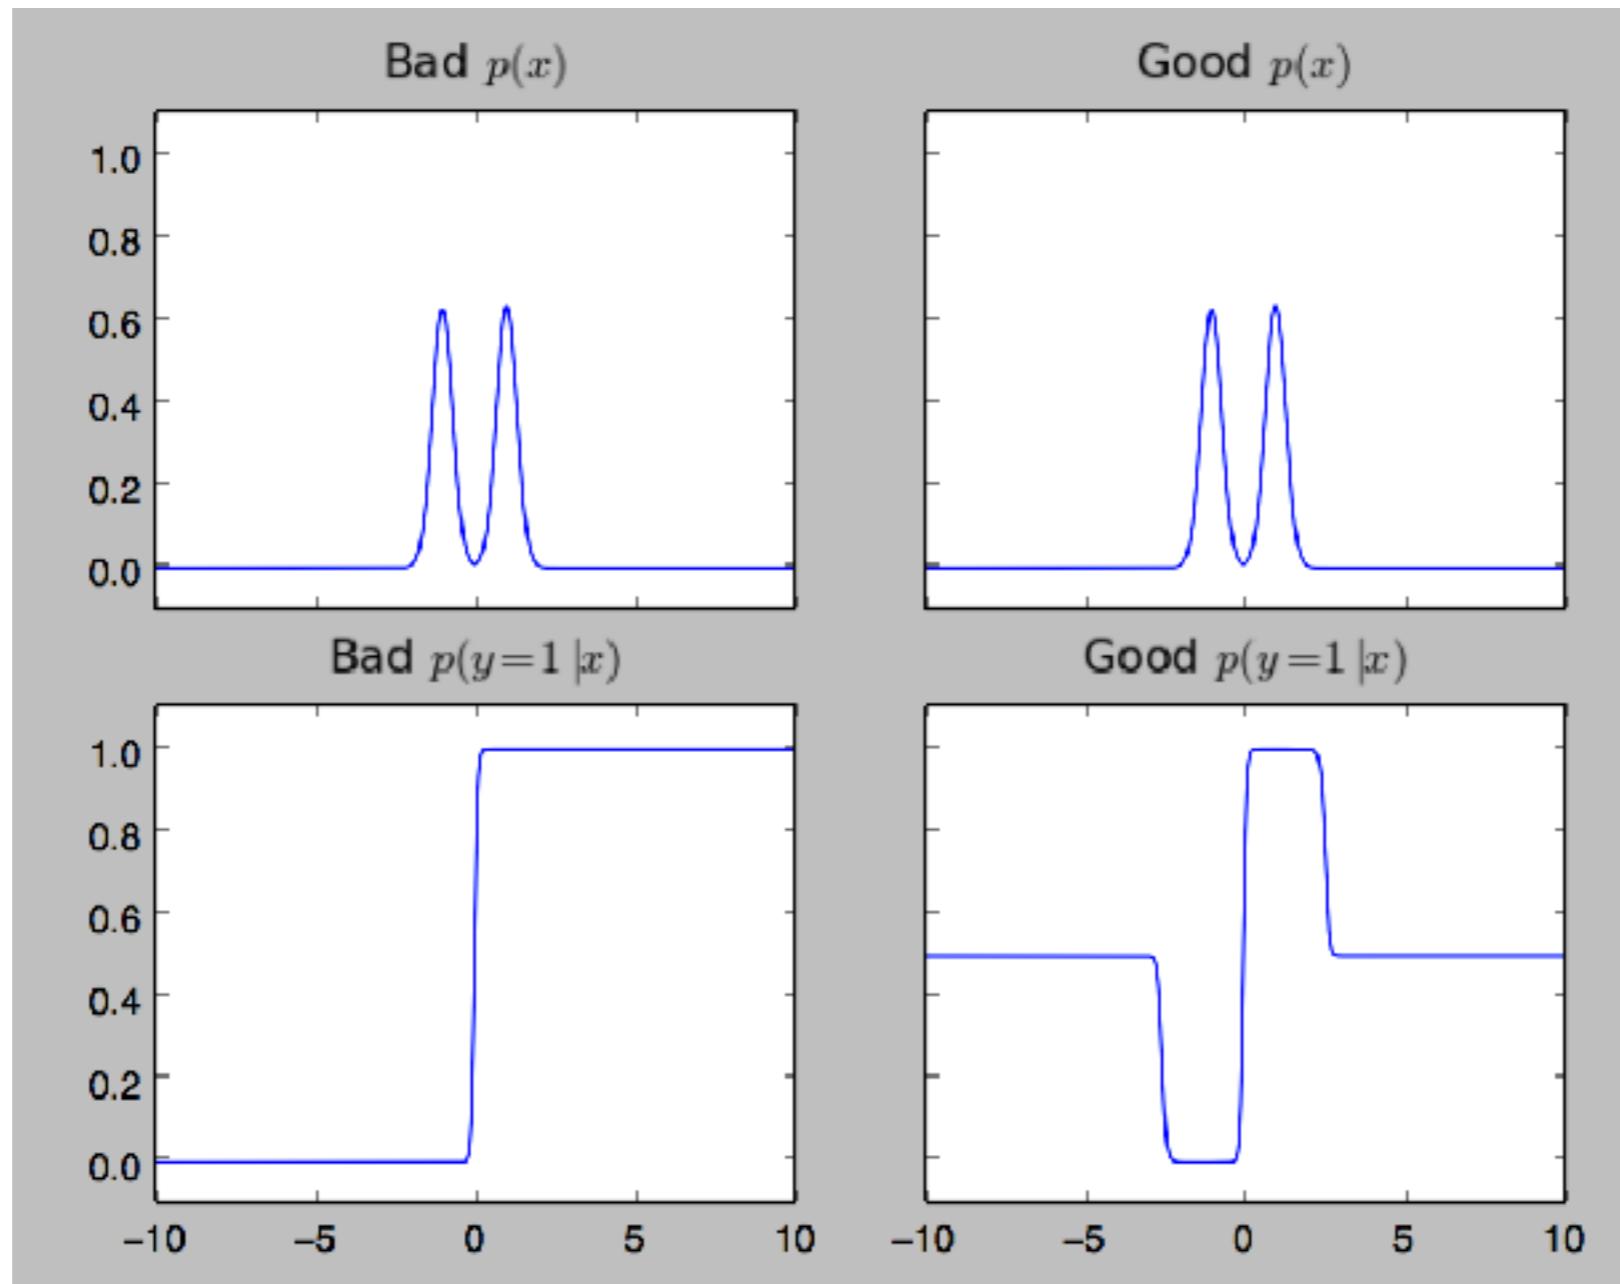
Removing perturbation
with an autoencoder
Ensembles

Error correcting
codes

Multiple glimpses

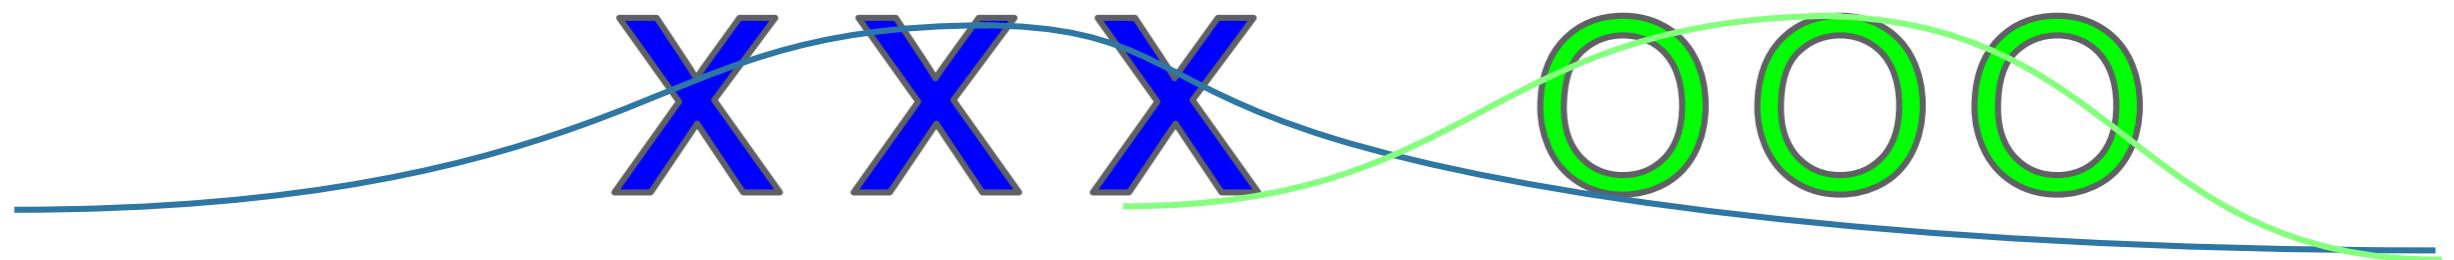
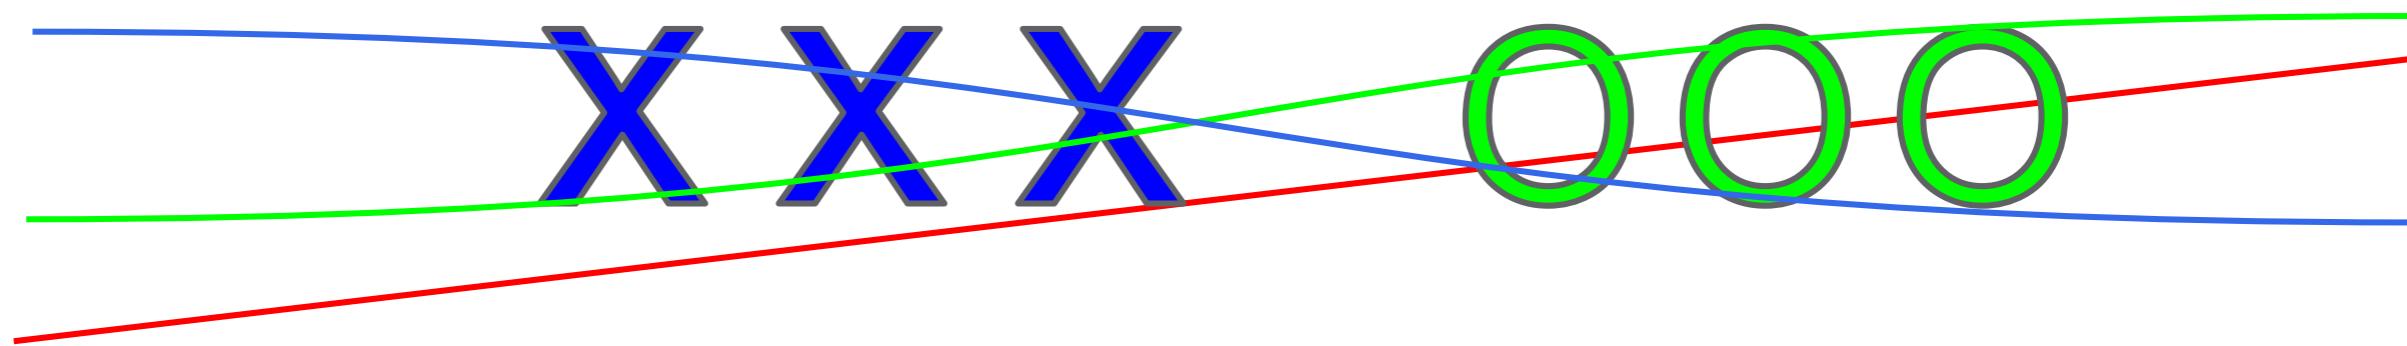
Adding noise
at train time

Generative Modeling is not Sufficient to Solve the Problem



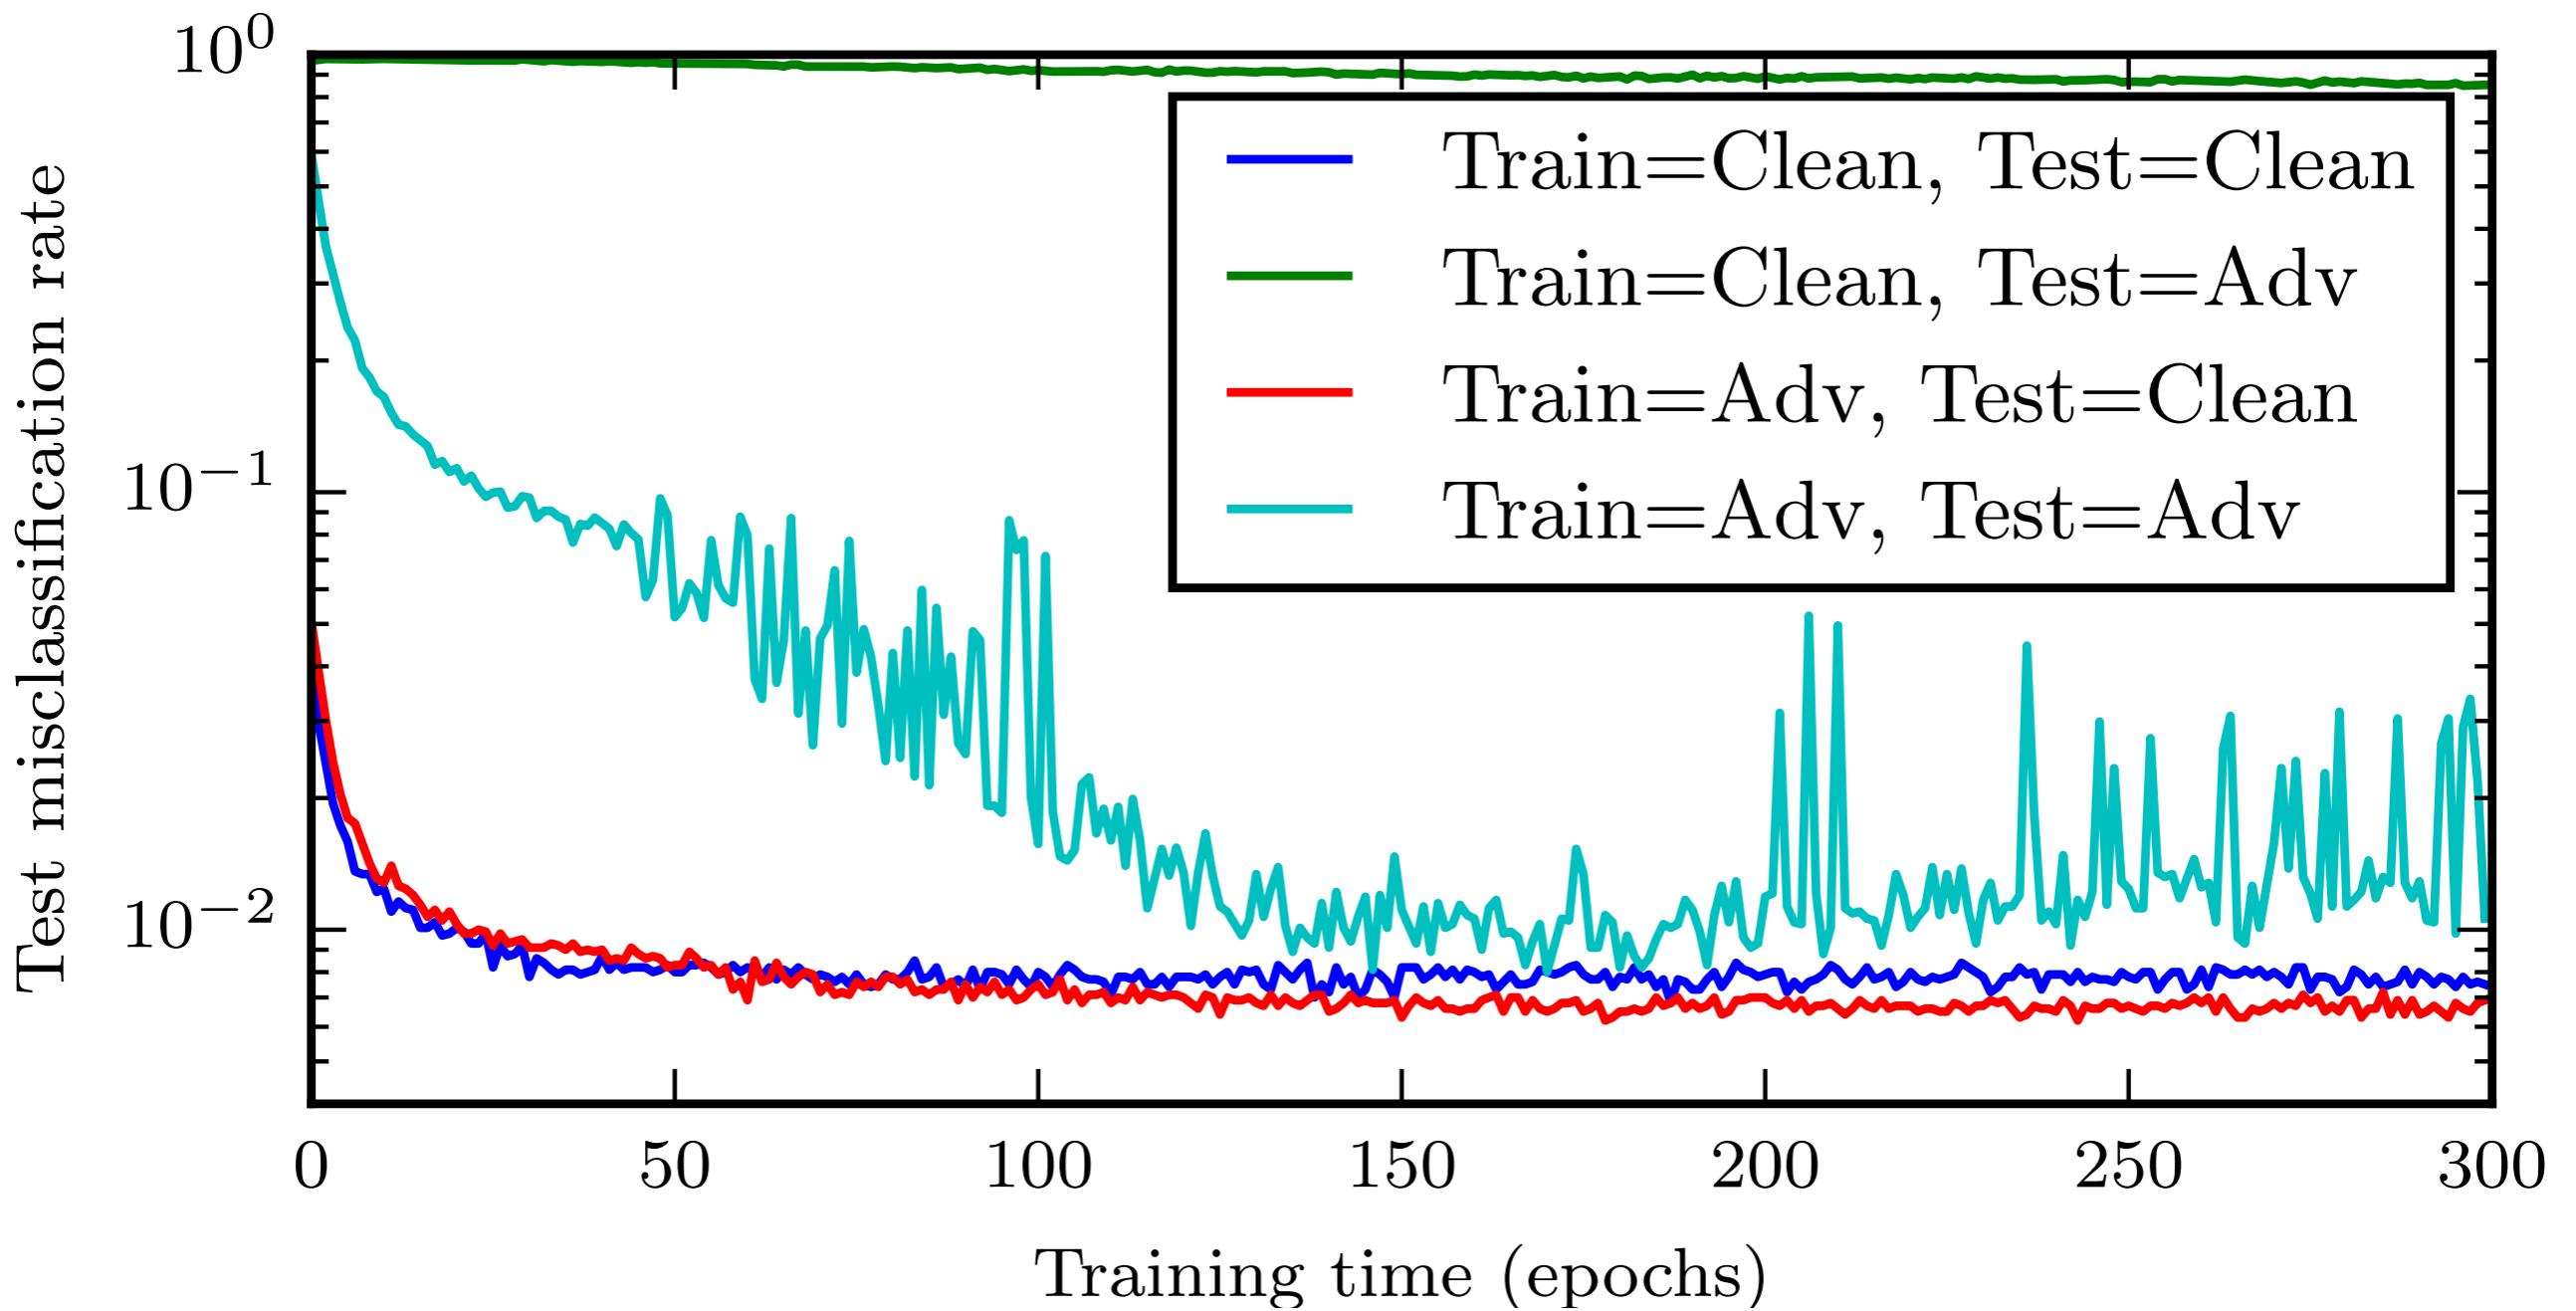
Universal approximator theorem

Neural nets can represent either function:



Maximum likelihood doesn't cause them to learn the right function. But we can fix that...

Training on Adversarial Examples



Adversarial Training of other Models

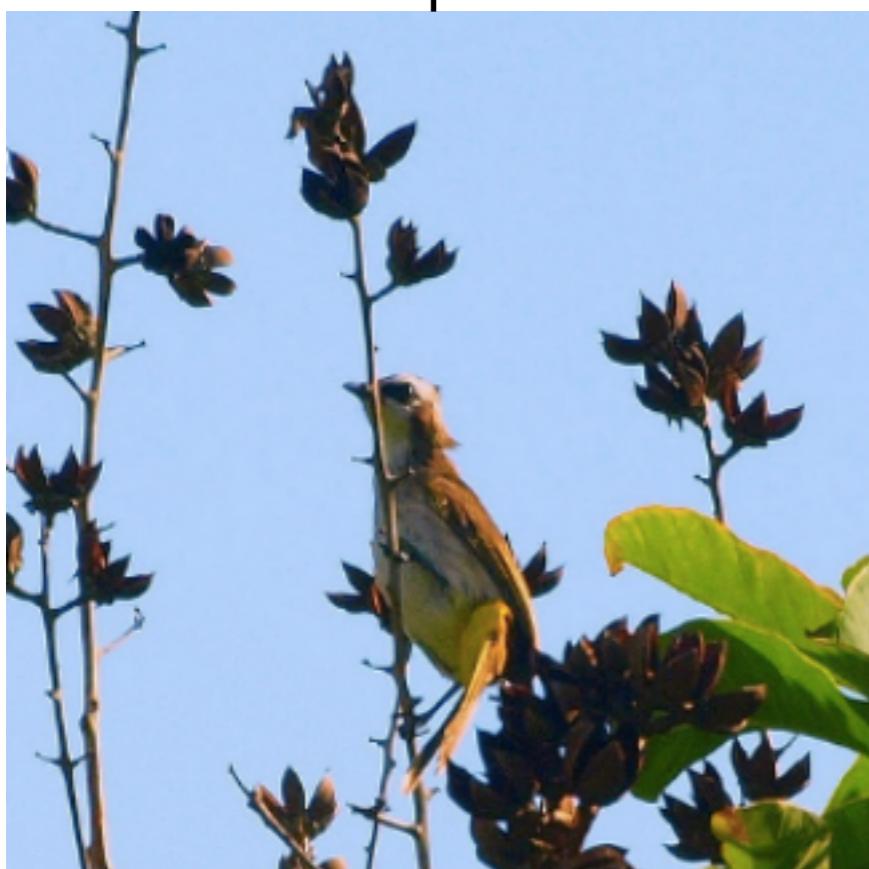
- Linear models: SVM / linear regression cannot learn a step function, so adversarial training is less useful, very similar to weight decay
- k -NN: adversarial training is prone to overfitting.
- Takeaway: neural nets can actually become more secure than other models. *Adversarially trained neural nets have the best empirical success rate on adversarial examples of any machine learning model.*

Weaknesses Persist



Adversarial Training

Labeled as bird



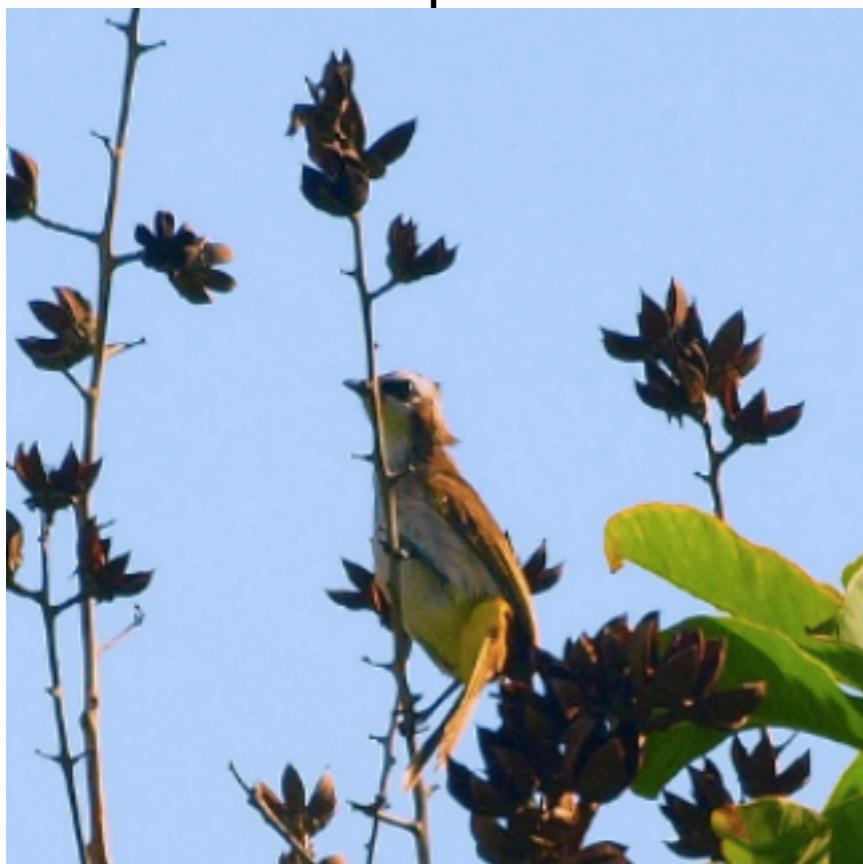
Still has same label (bird)



Decrease
probability
of bird class

Virtual Adversarial Training

Unlabeled; model
guesses it's probably
a bird, maybe a plane



New guess should
match old guess

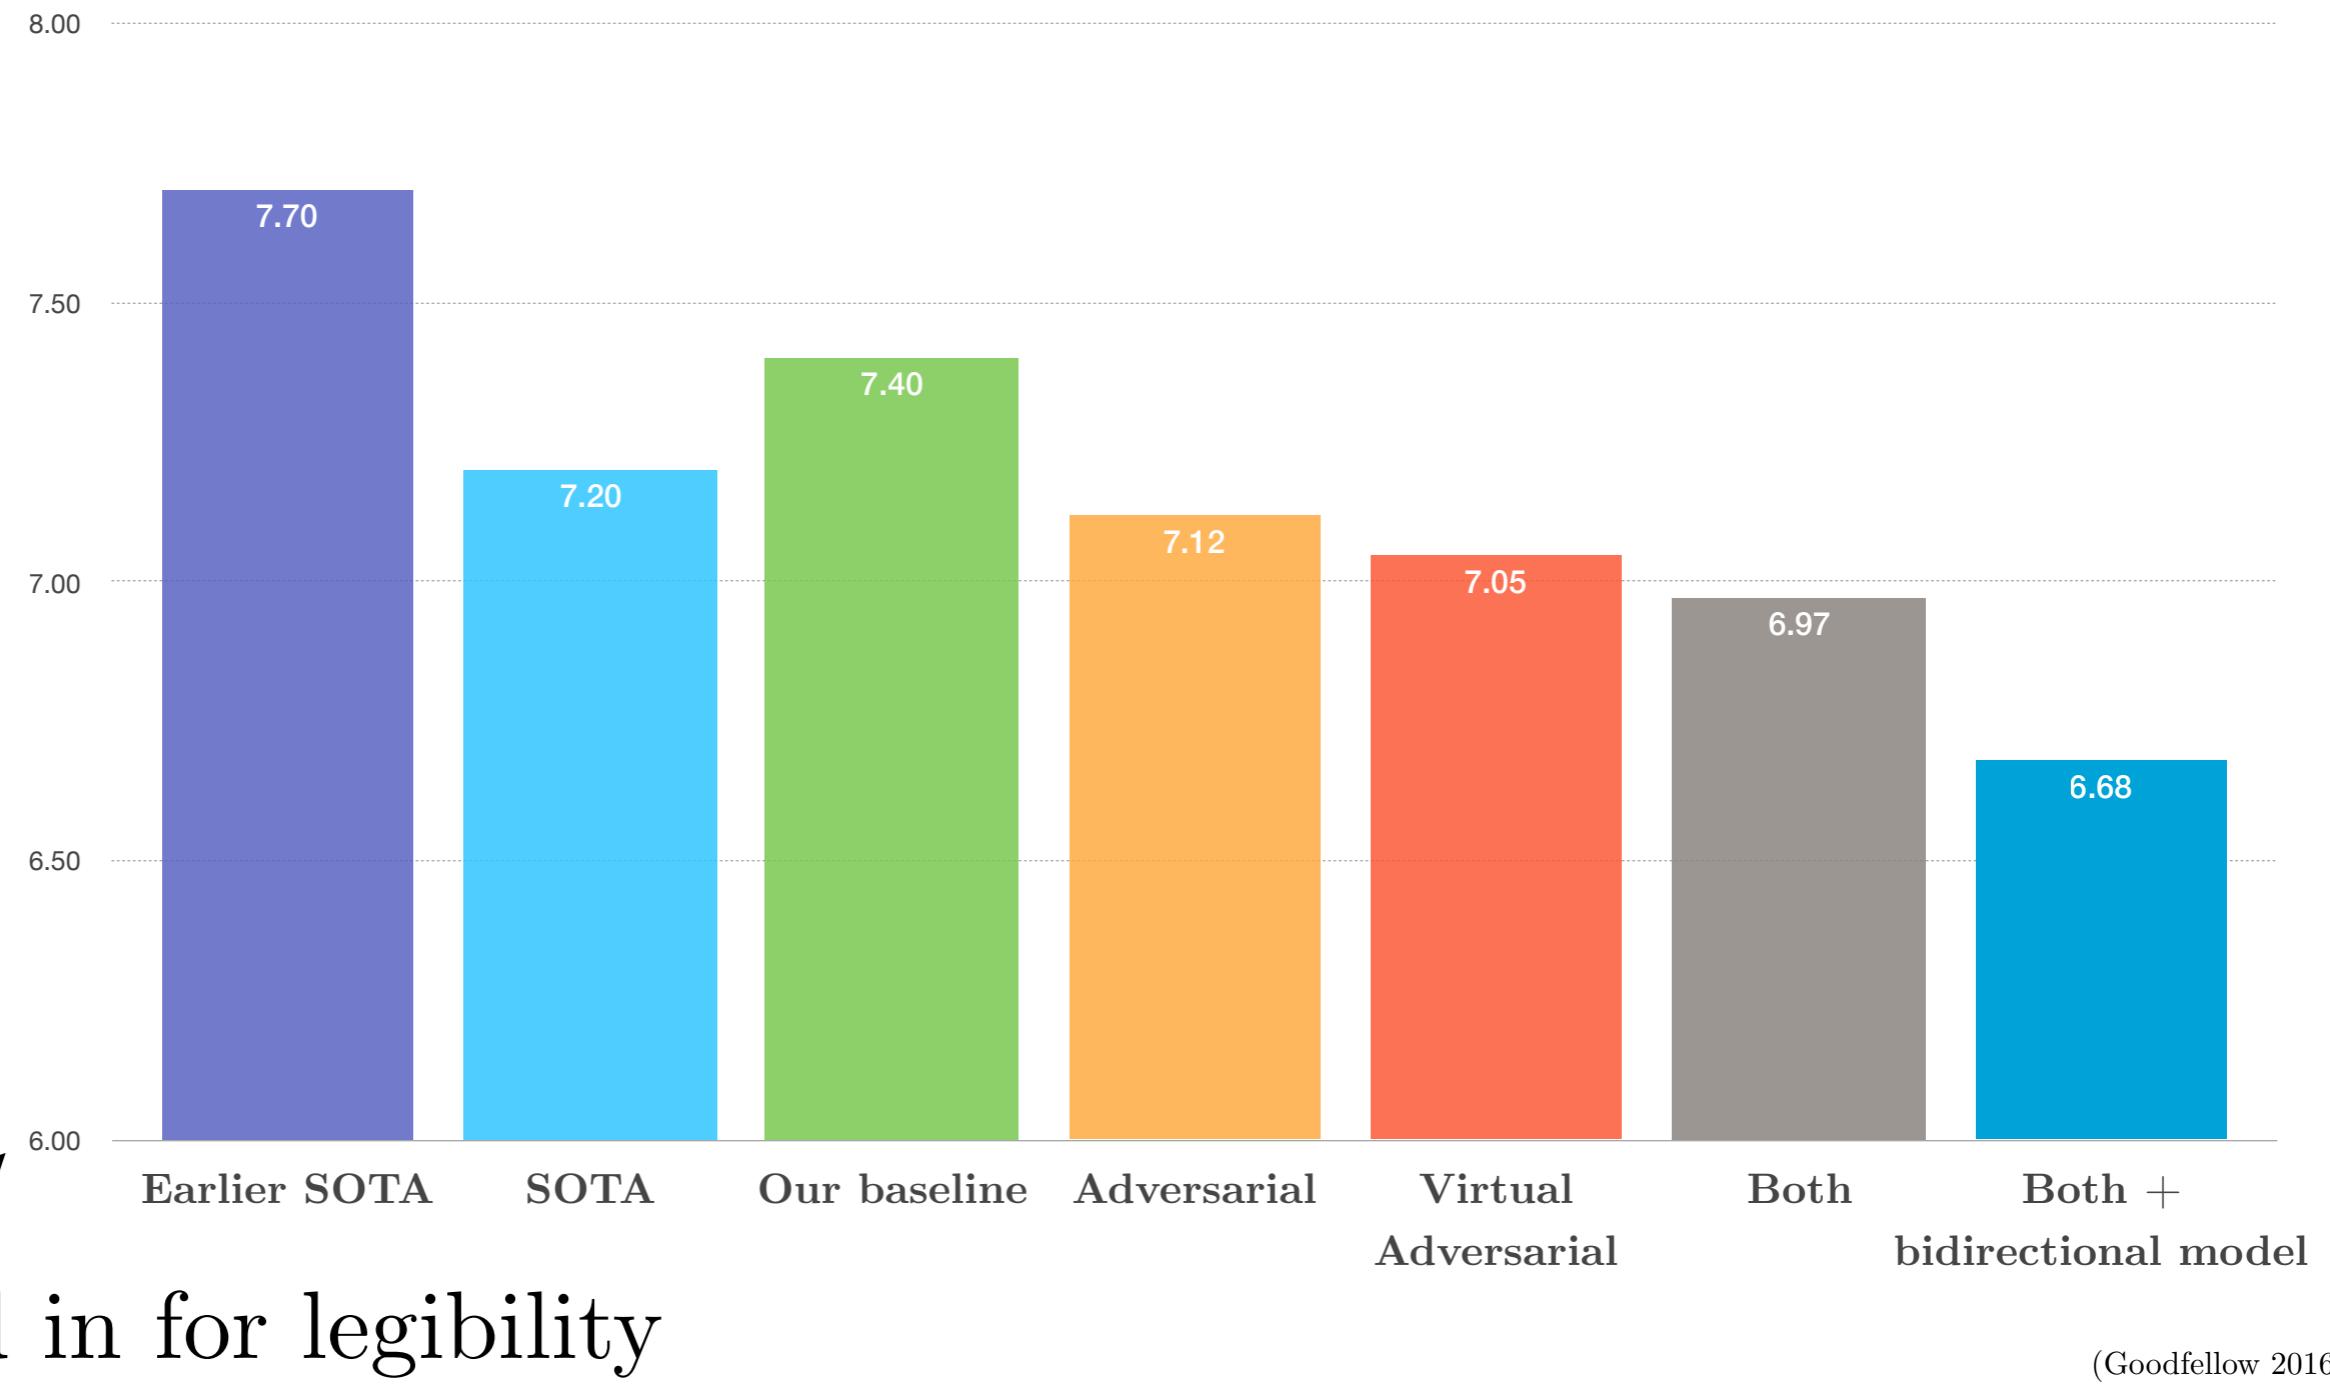
(probably bird, maybe plane)



→
Adversarial
perturbation
intended to
change the guess

Text Classification with VAT

RCV1 Misclassification Rate



Universal engineering machine (model-based optimization)

Make new inventions
by finding input
that maximizes
model's predicted
performance

Training data

Extrapolation



Conclusion

- Attacking is easy
- Defending is difficult
- Adversarial training provides regularization and semi-supervised learning
- The out-of-domain input problem is a bottleneck for model-based optimization generally

cleverhans

Open-source library available at:

<https://github.com/openai/cleverhans>

Built on top of TensorFlow (Theano support anticipated)

Standard implementation of attacks, for adversarial training
and reproducible benchmarks

